

Github Classroom

Repository : <https://github.com/cs418-fa24/project-check-in-team-7>

Project description

This project focuses on predicting the likelihood of diabetic patients being readmitted to the hospital within 30 days. By leveraging a dataset containing information about patient encounters, including demographic details, medical history, diagnosis codes, and treatment data, the goal is to develop a model that can identify high-risk patients. By predicting which patients are most likely to be readmitted, hospitals can allocate resources more effectively, prioritize interventions for high-risk individuals, and improve patient care.

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import OrdinalEncoder, OneHotEncoder, StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import LabelEncoder
from sklearn.impute import SimpleImputer
```


```
In [2]: df = pd.read_csv("diabetic_data_initial.csv")
```

```
In [3]: df.head()
```

```
Out[3]:
```

	encounter_id	patient_nbr	race	gender	age	weight	admission_type_id	dis
0	2278392	8222157	Caucasian	Female	[0-10)	?	6	
1	149190	55629189	Caucasian	Female	[10-20)	?	1	
2	64410	86047875	AfricanAmerican	Female	[20-30)	?	1	
3	500364	82442376	Caucasian	Male	[30-40)	?	1	
4	16680	42519267	Caucasian	Male	[40-50)	?	1	

5 rows × 50 columns



Data Exploration

```
In [4]: print(f"Dataset shape: {df.shape}")
```

Dataset shape: (101766, 50)

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 101766 entries, 0 to 101765
Data columns (total 50 columns):
```

#	Column	Non-Null Count	Dtype
0	encounter_id	101766 non-null	int64
1	patient_nbr	101766 non-null	int64
2	race	101766 non-null	object
3	gender	101766 non-null	object
4	age	101766 non-null	object
5	weight	101766 non-null	object
6	admission_type_id	101766 non-null	int64
7	discharge_disposition_id	101766 non-null	int64
8	admission_source_id	101766 non-null	int64
9	time_in_hospital	101766 non-null	int64
10	payer_code	101766 non-null	object
11	medical_specialty	101766 non-null	object
12	num_lab_procedures	101766 non-null	int64
13	num_procedures	101766 non-null	int64
14	num_medications	101766 non-null	int64
15	number_outpatient	101766 non-null	int64
16	number_emergency	101766 non-null	int64
17	number_inpatient	101766 non-null	int64
18	diag_1	101766 non-null	object
19	diag_2	101766 non-null	object
20	diag_3	101766 non-null	object
21	number_diagnoses	101766 non-null	int64
22	max_glu_serum	5346 non-null	object
23	A1Cresult	17018 non-null	object
24	metformin	101766 non-null	object
25	repaglinide	101766 non-null	object
26	nateglinide	101766 non-null	object
27	chlorpropamide	101766 non-null	object
28	glimepiride	101766 non-null	object
29	acetohexamide	101766 non-null	object
30	glipizide	101766 non-null	object
31	glyburide	101766 non-null	object
32	tolbutamide	101766 non-null	object
33	pioglitazone	101766 non-null	object
34	rosiglitazone	101766 non-null	object
35	acarbose	101766 non-null	object
36	miglitol	101766 non-null	object
37	troglitazone	101766 non-null	object
38	tolazamide	101766 non-null	object
39	examide	101766 non-null	object
40	citoglipton	101766 non-null	object
41	insulin	101766 non-null	object
42	glyburide-metformin	101766 non-null	object
43	glipizide-metformin	101766 non-null	object
44	glimepiride-pioglitazone	101766 non-null	object
45	metformin-rosiglitazone	101766 non-null	object
46	metformin-pioglitazone	101766 non-null	object
47	change	101766 non-null	object
48	diabetesMed	101766 non-null	object
49	readmitted	101766 non-null	object

dtypes: int64(13), object(37)
memory usage: 38.8+ MB

```
In [6]: # Display statistical summary for numerical columns
print(df.describe())
```

	encounter_id	patient_nbr	admission_type_id \
count	1.017660e+05	1.017660e+05	101766.000000
mean	1.652016e+08	5.433040e+07	2.024006
std	1.026403e+08	3.869636e+07	1.445403
min	1.252200e+04	1.350000e+02	1.000000
25%	8.496119e+07	2.341322e+07	1.000000
50%	1.523890e+08	4.550514e+07	1.000000
75%	2.302709e+08	8.754595e+07	3.000000
max	4.438672e+08	1.895026e+08	8.000000

	discharge_disposition_id	admission_source_id	time_in_hospital \
count	101766.000000	101766.000000	101766.000000
mean	3.715642	5.754437	4.395987
std	5.280166	4.064081	2.985108
min	1.000000	1.000000	1.000000
25%	1.000000	1.000000	2.000000
50%	1.000000	7.000000	4.000000
75%	4.000000	7.000000	6.000000
max	28.000000	25.000000	14.000000

	num_lab_procedures	num_procedures	num_medications	number_outpatient \
count	101766.000000	101766.000000	101766.000000	101766.000000
mean	43.095641	1.339730	16.021844	0.369357
std	19.674362	1.705807	8.127566	1.267265
min	1.000000	0.000000	1.000000	0.000000
25%	31.000000	0.000000	10.000000	0.000000
50%	44.000000	1.000000	15.000000	0.000000
75%	57.000000	2.000000	20.000000	0.000000
max	132.000000	6.000000	81.000000	42.000000

	number_emergency	number_inpatient	number_diagnoses
count	101766.000000	101766.000000	101766.000000
mean	0.197836	0.635566	7.422607
std	0.930472	1.262863	1.933600
min	0.000000	0.000000	1.000000
25%	0.000000	0.000000	6.000000
50%	0.000000	0.000000	8.000000
75%	0.000000	1.000000	9.000000
max	76.000000	21.000000	16.000000

```
In [7]: # Display unique values for categorical columns
object_columns = df.select_dtypes(include=['object'])
object_columns.head(10)
for column in object_columns:
    unique_values = df[column].unique()
    print(f"Unique values in '{column}': {unique_values}\n")
```

Unique values in 'race': ['Caucasian' 'AfricanAmerican' '?' 'Other' 'Asian' 'Hispanic']

Unique values in 'gender': ['Female' 'Male' 'Unknown/Invalid']

Unique values in 'age': ['[0-10)' '[10-20)' '[20-30)' '[30-40)' '[40-50)' '[50-60)' '[60-70)' '[70-80)' '[80-90)' '[90-100)']

Unique values in 'weight': ['?' '[75-100)' '[50-75)' '[0-25)' '[100-125)' '[25-50)' '[125-150)' '[175-200)' '[150-175)' '>200']

Unique values in 'payer_code': ['?' 'MC' 'MD' 'HM' 'UN' 'BC' 'SP' 'CP' 'SI' 'DM' 'CM' 'CH' 'PO' 'WC' 'OT' 'OG' 'MP' 'FR']

Unique values in 'medical_specialty': ['Pediatrics-Endocrinology' '?' 'InternalMedicine'

'Family/GeneralPractice' 'Cardiology' 'Surgery-General' 'Orthopedics' 'Gastroenterology' 'Surgery-Cardiovascular/Thoracic' 'Nephrology' 'Orthopedics-Reconstructive' 'Psychiatry' 'Emergency/Trauma' 'Pulmonology' 'Surgery-Neuro' 'Obstetrics&Gynecology-GynecologicOnco' 'ObstetricsandGynecology' 'Pediatrics' 'Hematology/Oncology' 'Otolaryngology' 'Surgery-Colon&Rectal' 'Pediatrics-CriticalCare' 'Endocrinology' 'Urology' 'Psychiatry-Child/Adolescent' 'Pediatrics-Pulmonology' 'Neurology' 'Anesthesiology-Pediatric' 'Radiology' 'Pediatrics-Hematology-Oncology' 'Psychology' 'Podiatry' 'Gynecology' 'Oncology' 'Pediatrics-Neurology' 'Surgery-Plastic' 'Surgery-Thoracic' 'Surgery-PlasticwithinHeadandNeck' 'Ophthalmology' 'Surgery-Pediatric' 'Pediatrics-EmergencyMedicine' 'PhysicalMedicineandRehabilitation' 'InfectiousDiseases' 'Anesthesiology' 'Rheumatology' 'AllergyandImmunology' 'Surgery-Maxillofacial' 'Pediatrics-InfectiousDiseases' 'Pediatrics-AllergyandImmunology' 'Dentistry' 'Surgeon' 'Surgery-Vascular' 'Osteopath' 'Psychiatry-Addictive' 'Surgery-Cardiovascular' 'PhysicianNotFound' 'Hematology' 'Proctology' 'Obstetrics' 'SurgicalSpecialty' 'Radiologist' 'Pathology' 'Dermatology' 'SportsMedicine' 'Speech' 'Hospitalist' 'OutreachServices' 'Cardiology-Pediatric' 'Perinatology' 'Neurophysiology' 'Endocrinology-Metabolism' 'DCPTEAM' 'Resident']

Unique values in 'diag_1': ['250.83' '276' '648' '8' '197' '414' '428' '398' '434' '250.7' '157'

'518' '999' '410' '682' '402' '737' '572' 'V57' '189' '786' '427' '996' '277' '584' '462' '473' '411' '174' '486' '998' '511' '432' '626' '295' '196' '250.6' '618' '182' '845' '423' '808' '250.4' '722' '403' '250.11' '784' '707' '440' '151' '715' '997' '198' '564' '812' '38' '590' '556' '578' '250.32' '433' 'V58' '569' '185' '536' '255' '250.13' '599' '558' '574' '491' '560' '244' '250.03' '577' '730' '188' '824' '250.8' '332' '562' '291' '296' '510' '401' '263' '438' '70' '250.02' '493' '642' '625' '571' '738' '593' '250.42' '807' '456' '446' '575' '250.41' '820' '515' '780' '250.22' '995' '235' '250.82' '721' '787' '162' '724' '282' '514' 'V55' '281' '250.33' '530' '466' '435' '250.12' 'V53' '789' '566' '822' '191' '557' '733' '455' '711' '482' '202' '280' '553' '225' '154' '441' '250.81' '349' '?' '962' '592' '507' '386' '156' '200' '728' '348' '459' '426' '388' '607' '337' '82' '531' '596' '288' '656' '573' '492' '220'

'516' '210' '922' '286' '885' '958' '661' '969' '250.93' '227' '112'
 '404' '823' '532' '416' '346' '535' '453' '250' '595' '211' '303'
 '250.01' '852' '218' '782' '540' '457' '285' '431' '340' '550' '54' '351'
 '601' '723' '555' '153' '443' '380' '204' '424' '241' '358' '694' '331'
 '345' '681' '447' '290' '158' '579' '436' '335' '309' '654' '805' '799'
 '292' '183' '78' '851' '458' '586' '311' '892' '305' '293' '415' '591'
 '794' '803' '79' '655' '429' '278' '658' '598' '729' '585' '444' '604'
 '727' '214' '552' '284' '680' '708' '41' '644' '481' '821' '413' '437'
 '968' '756' '632' '359' '275' '512' '781' '420' '368' '522' '294' '825'
 '135' '304' '320' '250.31' '669' '868' '496' '250.43' '826' '567' '3'
 '203' '53' '251' '565' '161' '495' '49' '250.1' '297' '663' '576' '355'
 '850' '287' '250.2' '611' '840' '350' '726' '537' '620' '180' '366' '783'
 '11' '751' '716' '250.3' '199' '464' '580' '836' '664' '283' '813' '966'
 '289' '965' '184' '480' '608' '333' '972' '212' '117' '788' '924' '959'
 '621' '238' '785' '714' '942' '250.23' '710' '47' '933' '508' '478' '844'
 '7' '736' '233' '42' '250.5' '397' '395' '201' '421' '253' '250.92' '600'
 '494' '977' '39' '659' '312' '614' '647' '652' '646' '274' '861' '425'
 '527' '451' '485' '217' '250.53' '442' '970' '193' '160' '322' '581'
 '475' '623' '374' '582' '568' '465' '801' '237' '376' '150' '461' '913'
 '226' '617' '987' '641' '298' '790' '336' '362' '228' '513' '383' '746'
 '353' '911' '506' '873' '155' '860' '534' '802' '141' 'V45' '396' '310'
 '341' '242' '719' '239' '533' '616' '519' '301' 'V66' '5' '989' '230'
 '385' '300' '853' '871' '570' '848' '463' '9' '934' '250.21' '236' '361'
 '594' '501' '810' '643' '430' '528' '205' '791' '983' '992' '490' '172'
 '171' '622' '306' '863' '864' '474' '660' '759' '356' '634' '967' '551'
 '695' '187' '732' '747' '323' '308' '370' '252' '152' '846' '164' '365'
 '718' '48' '266' '720' '94' '344' '797' '170' '878' '904' 'V56' '882'
 '843' '709' '973' '454' '686' '939' '487' '229' '991' '483' '357' '692'
 '796' '693' '935' '936' '800' '920' 'V26' '261' '307' '262' '250.9' '831'
 '145' '223' 'V71' '839' '685' 'V54' '35' '34' '179' '964' '136' '324'
 '389' '815' '334' '143' '526' '588' '192' 'V67' '394' '917' '88' '219'
 '325' '792' '717' '994' '990' '793' '207' '637' '195' '373' '847' '827'
 '31' '891' '814' 'V60' '703' '865' '352' '627' '378' '342' '886' '369'
 '745' '705' '816' '541' '986' '610' '633' '640' '753' '173' '835' '379'
 '445' '272' '382' '945' '619' '881' '250.52' '866' '405' '916' '215'
 '893' '75' '671' '928' '906' '897' '725' '867' '115' '890' '734' '521'
 '674' '470' '834' '146' '696' '524' '980' '691' '384' '142' '879'
 '250.51' '246' '208' '448' '955' '653' '149' '245' '735' '883' '854'
 '952' '838' '194' 'V43' '163' '216' '147' '354' '27' '477' '318' '880'
 '921' '377' '471' '683' '175' '602' '250.91' '982' '706' '375' '417'
 '131' '347' '870' '148' '862' '61' '817' '914' '360' '684' '314' 'V63'
 '36' '57' '240' '915' '971' '795' '988' '452' '963' '327' '731' '842'
 'V25' '645' '665' '110' '944' '603' '923' '412' '363' '957' '976' '698'
 '299' '700' '273' '974' '97' '529' '66' '98' '605' '941' '52' '806' '84'
 '271' '837' '657' '895' '338' '523' '542' '114' '543' '372' 'V70' 'E909'
 '583' 'V07' '422' '615' '279' '500' '903' '919' '875' '381' '804' '704'
 '23' '58' '649' '832' '133' '975' '833' '391' '690' '10' 'V51']

Unique values in 'diag_2': ['?' '250.01' '250' '250.43' '157' '411' '492' '427' '198' '403' '288'

'998' '507' '174' '425' '456' '401' '715' '496' '428' '585' '250.02'
 '410' '999' '996' '135' '244' '41' '571' '276' '997' '599' '424' '491'
 '553' '707' '286' '440' '493' '242' '70' 'V45' '250.03' '357' '511' '196'
 '396' '197' '414' '250.52' '577' '535' '413' '285' '53' '780' '518' '150'
 '566' '250.6' '867' '486' 'V15' '8' '788' '340' '574' '581' '228' '530'
 '250.82' '786' '294' '567' '785' '512' '305' '729' '250.51' '280' '648'

'560' '618' '444' '38' 'V10' '578' '277' '781' '250.42' '278' '426' '584'
 '462' '402' '153' '272' '733' '34' '881' '203' '250.41' '250.13' '293'
 '245' '250.12' '558' '787' '342' '573' '626' '303' '250.53' '458' '710'
 '415' 'V42' '284' '569' '759' '682' '112' '292' '435' '290' '250.93'
 '642' '536' '398' '319' '711' 'E878' '446' '255' 'V44' '250.7' '784'
 '300' '562' '162' '287' '447' '789' '790' '591' '200' '154' '304' '117'
 '847' '852' '250.83' '250.11' '816' '575' '416' '412' '441' '515' '372'
 '482' '382' 'V65' '572' '283' '78' '250.81' '576' '432' '595' '295' 'V12'
 '204' '466' '721' '434' '590' '271' '813' '368' '227' '783' '250.5' '258'
 '253' '309' '250.91' '519' '333' '459' '250.92' '250.4' '179' '420' '345'
 '433' '661' '537' '205' '722' '405' '437' '714' '211' 'E812' '263' '202'
 '397' '250.23' 'E932' '201' '301' '723' '614' '568' '861' 'V57' '724'
 '189' '297' '453' 'E888' '730' '354' '451' '738' 'E939' '805' 'V43' '155'
 '910' '218' '358' '220' 'E937' '583' '958' '794' '564' '436' '250.22'
 '620' '621' '331' '617' '596' '314' '378' '250.8' '625' '478' '731' '172'
 '404' '681' '470' '279' '281' '531' '443' '799' '337' '311' '719' 'E944'
 '423' 'E870' '465' 'E849' '782' '481' '480' 'V23' '199' '79' '438' '348'
 '42' 'E950' '473' '627' '726' '54' '490' '317' '332' '508' '369' '600'
 '349' '485' '208' '922' '431' '296' 'E934' '753' 'E935' '386' '728' '607'
 'E915' '344' '716' '289' '191' '873' '850' '611' '377' '352' '616' 'V17'
 '136' '455' '933' 'E885' '860' '513' '603' '484' '223' 'V72' '291' '151'
 'V58' '550' '510' '891' '185' '592' '791' '138' '598' '336' '362' '217'
 '825' '298' '821' 'E880' '343' '429' 'E879' '579' '225' '250.9' 'V49'
 '696' '233' '658' '969' '275' '250.1' '601' '704' '808' 'E890' 'V18'
 '920' '380' '570' 'E817' '359' '812' '274' 'V14' '324' '758' 'V66' '911'
 'E931' 'E924' '593' '792' '727' 'V46' '394' '532' 'V64' '557' '864' '718'
 'E942' '807' '604' '924' '820' '580' '273' '241' '282' '824' 'V61' '646'
 '701' '736' '565' '383' '250.2' 'E947' '452' '872' '905' 'E930' '921'
 '131' '448' '389' '421' '214' '705' '494' '752' '623' '9' '299' '959'
 '365' '967' 'E858' '40' '691' '909' '5' '814' '746' '250.31' '556' '680'
 '745' '351' '306' '110' '695' '552' '346' '918' '882' '947' '520' '188'
 '31' '356' '737' 'V08' '322' '182' '517' '974' 'E929' 'V53' '912' '252'
 '608' '516' 'E933' '94' '702' '923' '594' '647' '111' '934' '430' '487'
 '709' '796' '156' '977' '915' '756' '840' '341' '259' '693' '725' 'V62'
 '528' '683' '953' '457' '501' 'E900' 'V09' '522' '919' '461' '506' '193'
 '483' 'E936' '717' '802' '335' 'V54' '320' '945' '906' '239' '454' '826'
 '823' 'E941' '226' '795' '684' '844' '250.33' '308' '615' '588' '712'
 '663' '706' '833' '741' '713' '533' 'E884' '586' '555' '755' 'E928' '742'
 '869' '962' 'V11' '543' '373' '870' '913' '152' '810' '965' '907' '908'
 '995' '845' '474' '442' '751' '323' '472' '464' '686' '250.32' '540'
 '251' '811' '652' '659' '851' '422' '815' '307' '325' '463' '992' '692'
 '521' '917' 'E965' '524' '916' 'E813' '173' '238' '137' '514' '312' '837'
 '355' '980' '622' '475' '500' '754' '261' '801' '868' '968' '381' '11'
 '250.21' '694' '610' '734' 'E814' '310' '130' '246' '892' '846' '634'
 '75' 'E927' 'E905' '183' '379' 'E917' '163' 'E868' '495' '747' '989'
 'E854' '240' '832' '605' '602' '644' 'V16' '35' 'V70' '376' '266' 'E918'
 '619' '477' '656' '46' '883' '171' 'V13' '698' '842' 'E850' '800' '269'
 '664' 'E887' '952' '164' 'E881' '527' '685' '366' '836' '27' 'V63' '865'
 '793' '232' '990' '52' '831' '327' '542' '806' '972' '862' 'E829' 'E919'
 '944' 'E916' '963' '316' '645' '347' 'V85' '374' 'V02' '748' '256' '186'
 '866' '975' '96' '395' '262' 'E819' '654' '994' '318' 'E826' '879' '674'
 '641' '822' '145' '797' '353' 'E938' 'E816' '948' '987' '99' '192'
 '250.3' 'E906' '534' '115' 'E818' 'E980' '360' '338' '529' '871' '750'
 '212' '302' '955' '141' '88' 'V25' '215' '350' 'V50' 'V03' 'E853' 'E968'
 'E882' '140' '703' '991' '893' 'E821' '235' 'V69' '670' '195' 'V55' '388'
 '268' '894' '114' '260' '853' '7' '880' 'V86' '180' 'E945' '523' '863'

'649' '270' '665' '460' '942' '364' '66' 'E883' '123' '884' 'V60' '843' '927']

Unique values in 'diag_3': ['?' '255' 'V27' '403' '250' 'V45' '38' '486' '996' '197' '250.6' '427' '627' '414' '416' '714' '428' '582' 'V43' '250.01' '263' '250.42' '276' '482' '401' '250.41' '585' '781' '278' '998' '568' '682' '618' '250.02' '305' '707' '496' '599' '715' '424' '518' '553' '794' '411' 'V42' '531' '511' '490' '562' '250.8' '250.7' '250.52' '784' '491' '581' '420' '8' '724' '730' '789' '131' '250.82' '999' '41' '493' '250.03' '753' '786' '529' 'E888' '425' '595' '303' '560' '711' '492' '332' '296' '438' '362' '250.4' '654' '244' 'V70' '737' '625' '681' '250.51' '404' 'V10' '810' '280' '440' '785' '588' '569' '272' '997' '250.43' '918' '584' '54' '788' '426' '722' '250.92' '196' '461' '535' '787' '891' '284' '458' '648' '780' '182' '285' '593' '413' '664' '564' '201' '356' 'V15' '292' '782' '473' '455' 'E932' '357' '348' '294' '250.23' '459' 'E878' '437' '733' '507' '525' '250.53' '397' '572' '805' '453' '331' '736' '402' '591' '576' '465' '533' '703' '349' '315' '658' '608' '578' '716' '382' '300' '282' '571' '536' '596' '287' '644' 'V11' '558' 'E885' '162' '198' '218' '412' '396' 'V14' '570' '433' 'E934' '882' '288' '577' '443' '729' '836' '295' '799' '281' '304' '153' '410' '616' '250.83' '601' '291' '75' '512' '660' '250.5' '598' '337' '574' '653' 'V58' '311' '415' '386' '602' '790' '112' '873' '620' '436' '70' '155' '138' '663' '530' '710' '42' '342' '250.91' 'E884' '515' '307' '704' '728' '731' '583' '238' '441' '293' '573' '532' '290' '594' '319' '250.13' '250.12' '519' '346' '380' '135' '642' '698' '924' '905' 'E933' '555' '309' 'E879' '286' '565' '752' '580' '446' '444' '344' '252' '35' '813' '394' '301' '575' '258' 'V17' '802' '435' '746' 'V12' '709' '881' 'E935' '139' '250.81' '718' '365' '202' '334' '185' '398' 'V44' '517' 'E849' '614' '466' '626' '250.9' '368' '605' '883' '289' '478' '617' '429' '442' 'V25' '866' '610' '557' '959' 'E942' '94' '920' '345' '313' '379' '79' '516' '586' '821' '600' '242' '373' '592' 'V64' '487' '253' '706' 'E947' '117' '340' 'E950' '656' 'E949' '590' 'V09' '250.22' '934' '694' '203' '250.93' '995' '726' '923' '958' '275' 'E929' '211' 'V18' 'V66' '199' '665' '53' '279' '522' '791' '890' '456' 'E938' 'E816' '122' '721' 'V65' '136' '480' '423' 'E920' '793' '647' '537' '351' '845' '336' '274' '719' '945' '434' '494' '227' '157' '208' '174' 'V57' '812' '734' '150' 'V23' '447' '692' '228' 'V16' '756' '405' 'E928' '823' '552' '528' '389' '240' '454' '792' '366' 'E939' '907' '270' '310' '266' '387' 'E931' '783' '245' '607' '355' 'E930' '705' '372' '369' '611' '283' 'V46' '110' '867' 'E956' '251' '250.2' '820' '712' '695' '567' '343' '723' 'V08' '273' '623' '807' '451' '495' '701' '34' 'V53' '314' '472' 'E945' '11' '189' '534' '354' '333' 'V54' '277' '659' '708' '452' '655' '816' '670' '621' '246' '953' '865' 'E817' '646' '151' '378' '78' '298' '840' '641' '521' '745' '619' '912' '506' 'E904' '259' 'E870' 'E980' '383' '204' '696' '566' '727' '47' 'E943' '358' '191' '965' '921' '432' '27' 'E861' '758' '477' '524' '751' '652' '556' '188' '825' '919' '732' '908' '951' '962' '685' 'E850' 'E944' '527' '341' '693' '250.1' 'V49' '860' '323' 'V55' '579' '508' '969' '205' '462' 'E880' '680' '697' '826' '200' '457' '717' '738' '742' '735' '235' '308' '725' '241' '824' '464' '260' '917' '239' '661' '892' '261' 'E883' '943' '744' 'E936' '796' '318' '967' '350' '854' 'E905' '9' '741' 'E941' '170' '643' '317' '759' '909' 'V22' '831' '713' '180' '801' '360' '359' '501' '335' '250.11' '306' '811' '690' 'V02' '271' '214' '847' '543' 'V63' '906' '842' '686' '445' '808' '861' 'E852' '220' 'E887' 'E858' '915' '970' '256' '747' '395' '243' '815' '481' '5' 'E927' '297' '299' '851' '864' '922' '384' 'E876' '225' '158' 'E937' '871' '88' '966' 'E917' 'E812'

'V62' 'E924' '604' '233' 'E916' '377' '797' 'V72' '172' '7' '421' '852'
 'E819' '972' '916' '956' '3' 'E965' '173' '193' '154' '347' '862' '250.3'
 '987' '470' '262' 'E855' '161' '115' '179' '910' '312' '17' '460' '265'
 '66' '163' 'V60' '870' 'E906' '514' '944' '844' '417' '152' '183' '991'
 '216' '385' '164' '935' '510' '814' '485' '850' '250.21' 'E919' '872'
 '195' '431' '597' '933' '171' '884' '156' '868' '483' 'E815' '542' 'V61'
 '853' '374' 'E881' 'E882' 'E822' '192' '754' '327' '523' '500' 'V85'
 '992' '657' '684' '603' 'E826' '550' '913' '376' '755' '361' '186' '720'
 '250.31' '674' '911' 'E813' '226' '365.44' 'E818' '146' '955' 'E894'
 '475' 'V13' '880' '930' 'E915' '381' '132' '353' '795' '893' 'V01' 'E853'
 '863' '540' 'E828' '430' '800' 'E865' '148' 'E946' '822' '879' '848'
 'V86' 'V03' '338' '989' '388' 'E966' '111' 'E922' '123' '757' 'E901'
 '141' '268' 'E892' '649' '702' '948' '223' '484' 'E886' '838' '928' '236'
 '624' '837' 'E987' 'V07' '841' '622' 'E912' 'E955' '463' 'V06' 'E864'
 '217' '877' '391' 'E825' '952' '669' '875' 'E900' '215' '538' '980' '834'
 '448' '175' '49' '876' '230' '57' 'E854' '942' '14' '750' '370' '671'
 '971']

Unique values in 'max_glu_serum': [nan '>300' 'Norm' '>200']

Unique values in 'A1Cresult': [nan '>7' '>8' 'Norm']

Unique values in 'metformin': ['No' 'Steady' 'Up' 'Down']

Unique values in 'repaglinide': ['No' 'Up' 'Steady' 'Down']

Unique values in 'nateglinide': ['No' 'Steady' 'Down' 'Up']

Unique values in 'chlorpropamide': ['No' 'Steady' 'Down' 'Up']

Unique values in 'glimepiride': ['No' 'Steady' 'Down' 'Up']

Unique values in 'acetohexamide': ['No' 'Steady']

Unique values in 'glipizide': ['No' 'Steady' 'Up' 'Down']

Unique values in 'glyburide': ['No' 'Steady' 'Up' 'Down']

Unique values in 'tolbutamide': ['No' 'Steady']

Unique values in 'pioglitazone': ['No' 'Steady' 'Up' 'Down']

Unique values in 'rosiglitazone': ['No' 'Steady' 'Up' 'Down']

Unique values in 'acarbose': ['No' 'Steady' 'Up' 'Down']

Unique values in 'miglitol': ['No' 'Steady' 'Down' 'Up']

Unique values in 'troglitazone': ['No' 'Steady']

Unique values in 'tolazamide': ['No' 'Steady' 'Up']

Unique values in 'examide': ['No']

Unique values in 'citoglipton': ['No']

Unique values in 'insulin': ['No' 'Up' 'Steady' 'Down']

Unique values in 'glyburide-metformin': ['No' 'Steady' 'Down' 'Up']

Unique values in 'glipizide-metformin': ['No' 'Steady']

Unique values in 'glimepiride-pioglitazone': ['No' 'Steady']

Unique values in 'metformin-rosiglitazone': ['No' 'Steady']

Unique values in 'metformin-pioglitazone': ['No' 'Steady']

Unique values in 'change': ['No' 'Ch']

Unique values in 'diabetesMed': ['No' 'Yes']

Unique values in 'readmitted': ['NO' '>30' '<30']

Data Preprocessing

```
In [8]: df.replace('?', np.nan, inplace=True)
```

```
In [9]: missing_values = df.isnull().sum()
print(missing_values[missing_values > 0])
```

```
race                2273
weight             98569
payer_code         40256
medical_specialty  49949
diag_1              21
diag_2              358
diag_3             1423
max_glu_serum      96420
A1Cresult          84748
dtype: int64
```

```
In [10]: # Drop columns with multiple missing values
df = df.drop(columns=['weight', 'payer_code', 'medical_specialty', 'discharge_dispos
```

```
In [11]: # Drop rows where the race column has missing values
df = df.dropna(subset=['race'])
```

```
In [12]: # Drop columns with only one unique value
df = df.drop(columns=['examide', 'citoglipton', 'metformin-rosiglitazone'])
```

```
In [13]: df['max_glu_serum'].unique()
```

```
Out[13]: array([nan, '>300', 'Norm', '>200'], dtype=object)
```

```
In [14]: # Replace the values in the 'max_glu_serum' column with more descriptive categories
df['max_glu_serum'] = df['max_glu_serum'].replace({
    '>300': 'Very High',
    '>200': 'High',
```

```

    'Norm': 'Normal',
    np.nan: 'No test performed'
})

print(df['max_glu_serum'].unique())

```

```
['No test performed' 'Very High' 'Normal' 'High']
```

```
In [15]: df['A1Cresult'].unique()
```

```
Out[15]: array([nan, '>7', '>8', 'Norm'], dtype=object)
```

```
In [16]: # Replace the values in the 'A1Cresult' column with more descriptive categories
df['A1Cresult'] = df['A1Cresult'].replace({
    '>8': 'Very High',
    '>7': 'High',
    'Norm': 'Normal',
    np.nan: 'No test performed'
})

print(df['A1Cresult'].unique())

```

```
['No test performed' 'High' 'Very High' 'Normal']
```

```
In [17]: # Mapping of diagnosis categories with corresponding ICD-9 code ranges
categories = {
    'Circulatory': ['390-459', '785'],
    'Respiratory': ['460-519', '786'],
    'Digestive': ['520-579', '787'],
    'Diabetes': ['250.xx'],
    'Injury': ['800-999'],
    'Musculoskeletal': ['710-739'],
    'Genitourinary': ['580-629', '788'],
    'Neoplasms': ['140-239', '780', '781', '784', '790-799', '240-249', '251-279', '68
    'Other': [
        '290-319', 'E', 'V', '280-289', '320-359', '630-679', '360-389',
        '740-759'
    ]
}

def categorize_diagnosis(code):
    if pd.isna(code):
        return 'Unknown'
    code = str(code)

    for category, ranges in categories.items():
        for icd_range in ranges:
            if '-' in icd_range:
                start, end = icd_range.split('-')
                if start <= code <= end:
                    return category
            elif icd_range.endswith('xx') and code.startswith(icd_range[:3]):
                return category
            elif code.startswith(icd_range):
                return category
    return 'Other'

```

```
df['primary_diagnosis'] = df['diag_1'].apply(categorize_diagnosis)
df['secondary_diagnosis'] = df['diag_2'].apply(categorize_diagnosis)
df['additional_diagnosis'] = df['diag_3'].apply(categorize_diagnosis)

df[['diag_1', 'primary_diagnosis', 'diag_2', 'secondary_diagnosis', 'diag_3', 'addi
```

Out[17]:

	diag_1	primary_diagnosis	diag_2	secondary_diagnosis	diag_3	additional_diagnosis
0	250.83	Diabetes	NaN	Unknown	NaN	Unknown
1	276	Neoplasms	250.01	Diabetes	255	Neoplasms
2	648	Other	250	Diabetes	V27	Other
3	8	Other	250.43	Diabetes	403	Circulatory
4	197	Neoplasms	157	Neoplasms	250	Diabetes

In [18]: `df = df.drop(columns=['diag_1', 'diag_2', 'diag_3'])`

In [19]: *# Map the numerical 'admission_type_id' to descriptive labels*

```
admission_type_mapping = {
    1: 'Emergency',
    2: 'Urgent',
    3: 'Elective',
    4: 'Newborn',
    5: 'Not Available',
    6: 'NULL',
    7: 'Trauma Center',
    8: 'Not Mapped'
}

df['admission_type'] = df['admission_type_id'].replace(admission_type_mapping)
df = df.drop(columns=['admission_type_id'])

print(df[['admission_type']].head())
```

```
admission_type
0      NULL
1  Emergency
2  Emergency
3  Emergency
4  Emergency
```

In [20]: `df.info()`

```

<class 'pandas.core.frame.DataFrame'>
Index: 99493 entries, 0 to 101765
Data columns (total 42 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   encounter_id                          99493 non-null  int64
1   patient_nbr                           99493 non-null  int64
2   race                                   99493 non-null  object
3   gender                                 99493 non-null  object
4   age                                    99493 non-null  object
5   time_in_hospital                      99493 non-null  int64
6   num_lab_procedures                    99493 non-null  int64
7   num_procedures                        99493 non-null  int64
8   num_medications                       99493 non-null  int64
9   number_outpatient                     99493 non-null  int64
10  number_emergency                       99493 non-null  int64
11  number_inpatient                       99493 non-null  int64
12  number_diagnoses                       99493 non-null  int64
13  max_glu_serum                          99493 non-null  object
14  A1Cresult                              99493 non-null  object
15  metformin                              99493 non-null  object
16  repaglinide                            99493 non-null  object
17  nateglinide                            99493 non-null  object
18  chlorpropamide                         99493 non-null  object
19  glimepiride                            99493 non-null  object
20  acetohexamide                         99493 non-null  object
21  glipizide                              99493 non-null  object
22  glyburide                              99493 non-null  object
23  tolbutamide                            99493 non-null  object
24  pioglitazone                           99493 non-null  object
25  rosiglitazone                          99493 non-null  object
26  acarbose                               99493 non-null  object
27  miglitol                               99493 non-null  object
28  troglitazone                           99493 non-null  object
29  tolazamide                             99493 non-null  object
30  insulin                                99493 non-null  object
31  glyburide-metformin                    99493 non-null  object
32  glipizide-metformin                    99493 non-null  object
33  glimepiride-pioglitazone                99493 non-null  object
34  metformin-pioglitazone                  99493 non-null  object
35  change                                 99493 non-null  object
36  diabetesMed                             99493 non-null  object
37  readmitted                             99493 non-null  object
38  primary_diagnosis                       99493 non-null  object
39  secondary_diagnosis                     99493 non-null  object
40  additional_diagnosis                    99493 non-null  object
41  admission_type                          99493 non-null  object
dtypes: int64(10), object(32)
memory usage: 32.6+ MB

```

Visualizations:

The plots below showcase how the different columns like Race, Gender, Age Groups are spread in our dataset

```
In [21]: import matplotlib.pyplot as plt
import seaborn as sns

# Race distribution
race_counts = df['race'].value_counts()
print(race_counts)

# Visualization: Race distribution
plt.figure(figsize=(8, 6))
sns.barplot(x=race_counts.index, y=race_counts.values, palette='Set2')
plt.title('Distribution of Race')
plt.xlabel('Race')
plt.ylabel('Frequency')
plt.show()

# Gender distribution
gender_counts = df['gender'].value_counts()
print(gender_counts)

# Visualization: Gender distribution
plt.figure(figsize=(8, 6))
sns.barplot(x=gender_counts.index, y=gender_counts.values, palette='Set1')
plt.title('Distribution of Gender')
plt.xlabel('Gender')
plt.ylabel('Frequency')
plt.show()

# Age distribution
age_counts = df['age'].value_counts()
print(age_counts)


# Visualization: Age distribution
plt.figure(figsize=(10, 6))
sns.barplot(x=age_counts.index, y=age_counts.values, palette='muted')
plt.title('Distribution of Age Groups')
plt.xlabel('Age Group')
plt.ylabel('Frequency')
plt.xticks(rotation=45)
plt.show()
```

```
race
Caucasian      76099
AfricanAmerican 19210
Hispanic       2037
Other          1506
Asian          641
Name: count, dtype: int64
```

<ipython-input-21-496251125ede>:10: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=race_counts.index, y=race_counts.values, palette='Set2')
```




```
gender
Female      53575
Male        45917
Unknown/Invalid  1
Name: count, dtype: int64
```

```
<ipython-input-21-496251125ede>:22: FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.
```

```
sns.barplot(x=gender_counts.index, y=gender_counts.values, palette='Set1')
```




```
age
[70-80)    25469
[60-70)    21988
[50-60)    16895
[80-90)    16800
[40-50)     9465
[30-40)     3699
[90-100)    2724
[20-30)     1611
[10-20)      682
[0-10)       160
```

Name: count, dtype: int64

<ipython-input-21-496251125ede>:34: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=age_counts.index, y=age_counts.values, palette='muted')
```

The plot presents the distribution of the number of lab procedures, procedures, and medications for a group of patients. The histograms show that the distribution of lab procedures and medications is right-skewed, meaning there are a few patients with many procedures and medications. At the same time, the majority have a lower number. The number of procedures is more evenly distributed, with a peak around 1. This suggests that most patients undergo only one procedure during their hospital stay.

```
In [22]: # Select the features of interest
features = ['num_lab_procedures', 'num_procedures', 'num_medications']

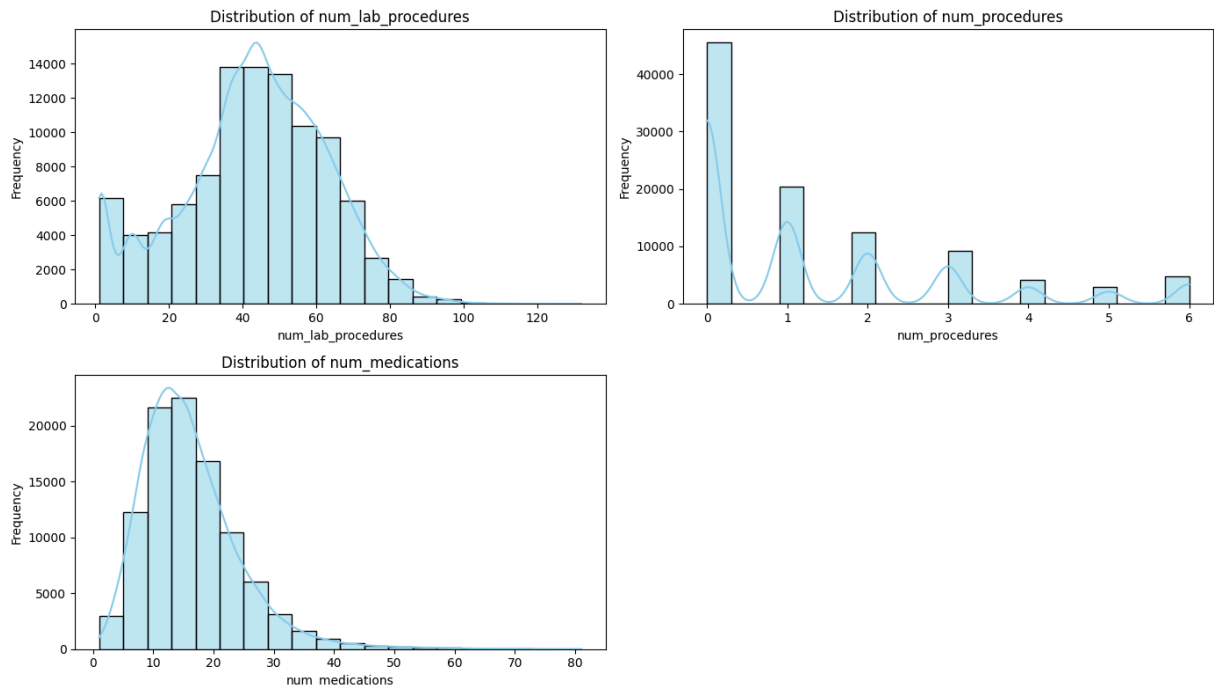
# Summary Statistics
summary_stats = df[features].describe(percentiles=[.25, .5, .75])
print("Summary Statistics for Lab and Procedure Counts:")
print(summary_stats)

# Visualization: Histograms for each feature
plt.figure(figsize=(14, 8))
for i, feature in enumerate(features, 1):
    plt.subplot(2, 2, i)
    sns.histplot(df[feature], bins=20, kde=True, color='skyblue')
    plt.title(f'Distribution of {feature}')
    plt.xlabel(feature)
    plt.ylabel('Frequency')

plt.tight_layout()
plt.show()
```

Summary Statistics for Lab and Procedure Counts:

	num_lab_procedures	num_procedures	num_medications
count	99493.000000	99493.000000	99493.000000
mean	43.072588	1.340577	16.026605
std	19.695858	1.703717	8.119790
min	1.000000	0.000000	1.000000
25%	31.000000	0.000000	10.000000
50%	44.000000	1.000000	15.000000
75%	57.000000	2.000000	20.000000
max	132.000000	6.000000	81.000000



The plots illustrate the frequency of medication changes and diabetes medication prescriptions within a dataset. The left plot shows that a majority of patients experienced medication changes (1 represents a change, 0 represents no change).

The right plot indicates that a significant proportion of patients were prescribed diabetes medication (1 represents a prescription, 0 represents no prescription). Both plots highlight the prevalence of medication adjustments and diabetes management within the studied population.

The plot presents the frequency of medication changes based on diabetes medication status. It shows that patients who were prescribed diabetes medication (indicated by 1) were more likely to experience medication changes compared to those who were not prescribed diabetes medication (indicated by 0). This suggests that diabetes medication use is associated with a higher frequency of medication adjustments, potentially due to the need for careful management of blood sugar levels and other related factors.

```
In [23]: # Frequency counts for 'change' and 'diabetesMed'
change_counts = df['change'].value_counts()
diabetesMed_counts = df['diabetesMed'].value_counts()

# Display frequency counts
```

```

print("Frequency of Medication Changes (change):")
print(change_counts)
print("\nFrequency of Diabetes Medication (diabetesMed):")
print(diabetesMed_counts)

# Visualization: Bar charts for 'change' and 'diabetesMed'
plt.figure(figsize=(12, 5))

# Bar chart for 'change'
plt.subplot(1, 2, 1)
sns.barplot(x=change_counts.index, y=change_counts.values, palette='Set1')
plt.title('Frequency of Medication Change')
plt.xlabel('Medication Change')
plt.ylabel('Frequency')

# Bar chart for 'diabetesMed'
plt.subplot(1, 2, 2)
sns.barplot(x=diabetesMed_counts.index, y=diabetesMed_counts.values, palette='Set2')
plt.title('Frequency of Diabetes Medication Prescription')
plt.xlabel('Diabetes Medication')
plt.ylabel('Frequency')

plt.tight_layout()
plt.show()

# Cross-tabulation of 'change' and 'diabetesMed'
change_diabetesMed_ct = pd.crosstab(df['change'], df['diabetesMed'])

# Display cross-tabulation
print("\nCross-tabulation of Change and Diabetes Medication:")
print(change_diabetesMed_ct)

# Visualization: Side-by-side bar chart for cross-tabulation
change_diabetesMed_ct.plot(kind='bar', color=['skyblue', 'salmon'], figsize=(8, 6))
plt.title('Medication Changes by Diabetes Medication Status')
plt.xlabel('Medication Change')
plt.ylabel('Frequency')
plt.legend(title='Diabetes Medication')
plt.tight_layout()
plt.show()

```

Frequency of Medication Changes (change):

change

No 53582

Ch 45911

Name: count, dtype: int64

Frequency of Diabetes Medication (diabetesMed):

diabetesMed

Yes 76492

No 23001

Name: count, dtype: int64

```
<ipython-input-23-cb4b548372a9>:16: FutureWarning:
```

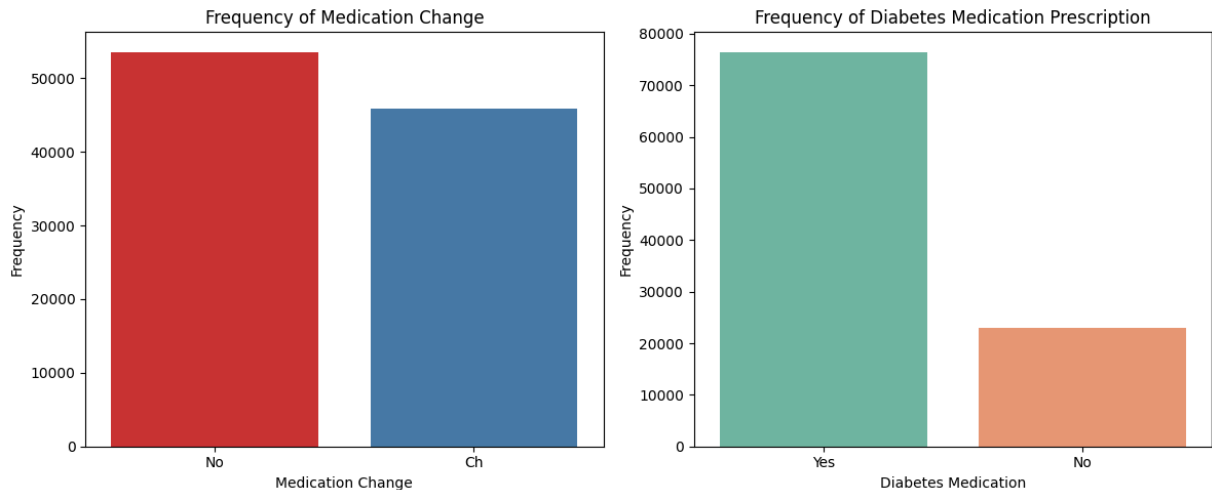
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=change_counts.index, y=change_counts.values, palette='Set1')
```

```
<ipython-input-23-cb4b548372a9>:23: FutureWarning:
```

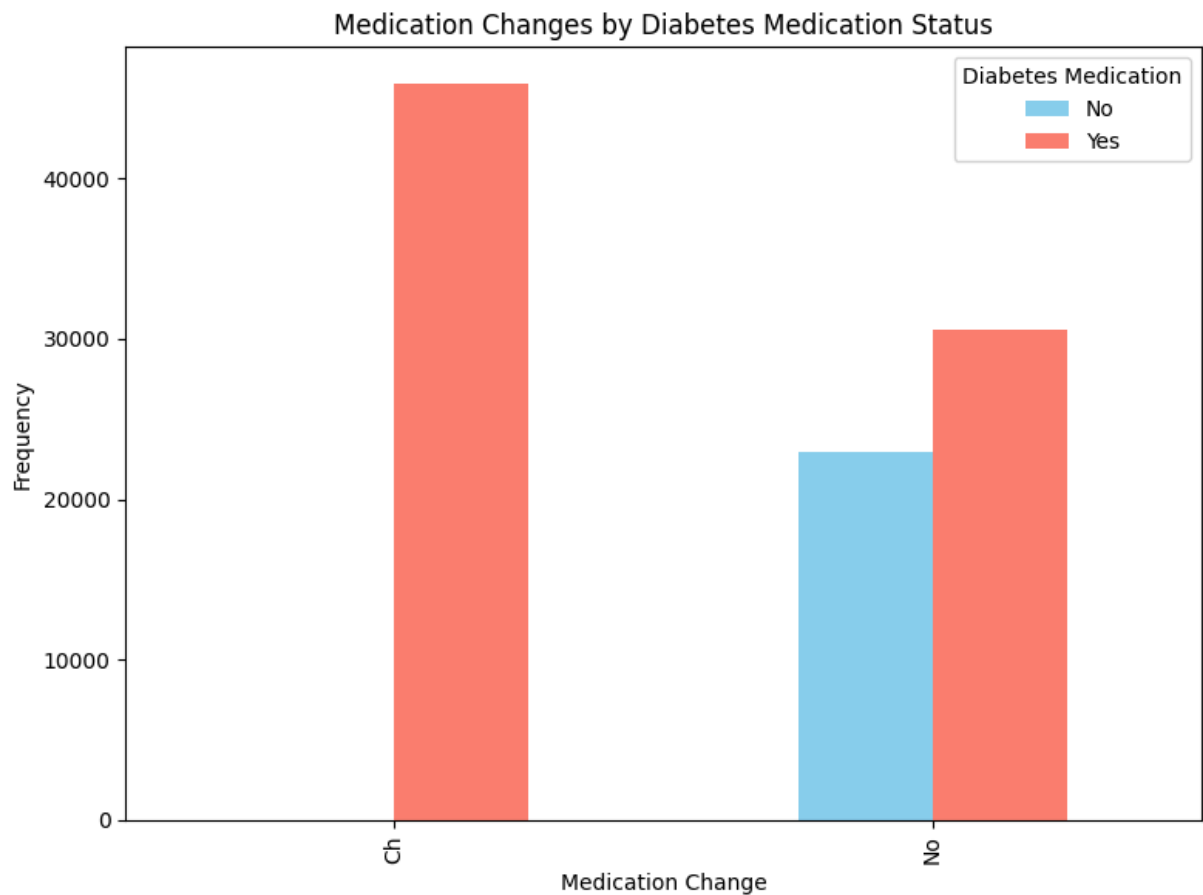
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=diabetesMed_counts.index, y=diabetesMed_counts.values, palette='Set 2')
```



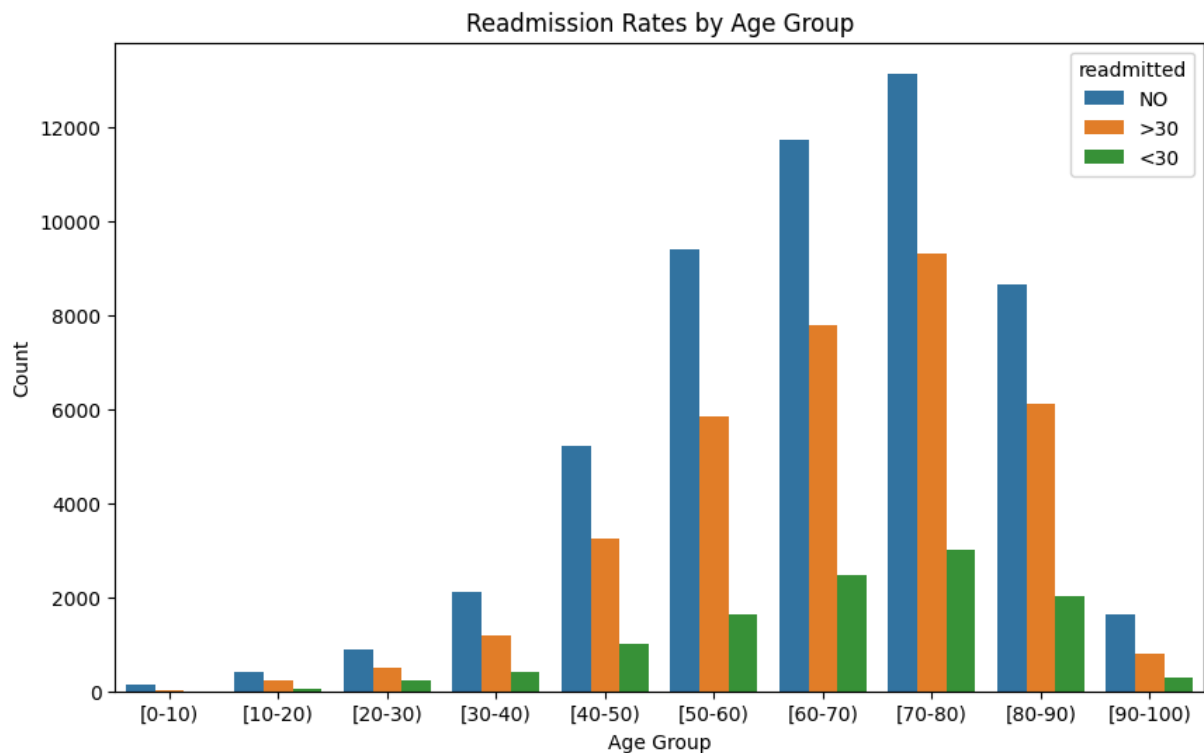
Cross-tabulation of Change and Diabetes Medication:

diabetesMed	No	Yes
change		
Ch	0	45911
No	23001	30581



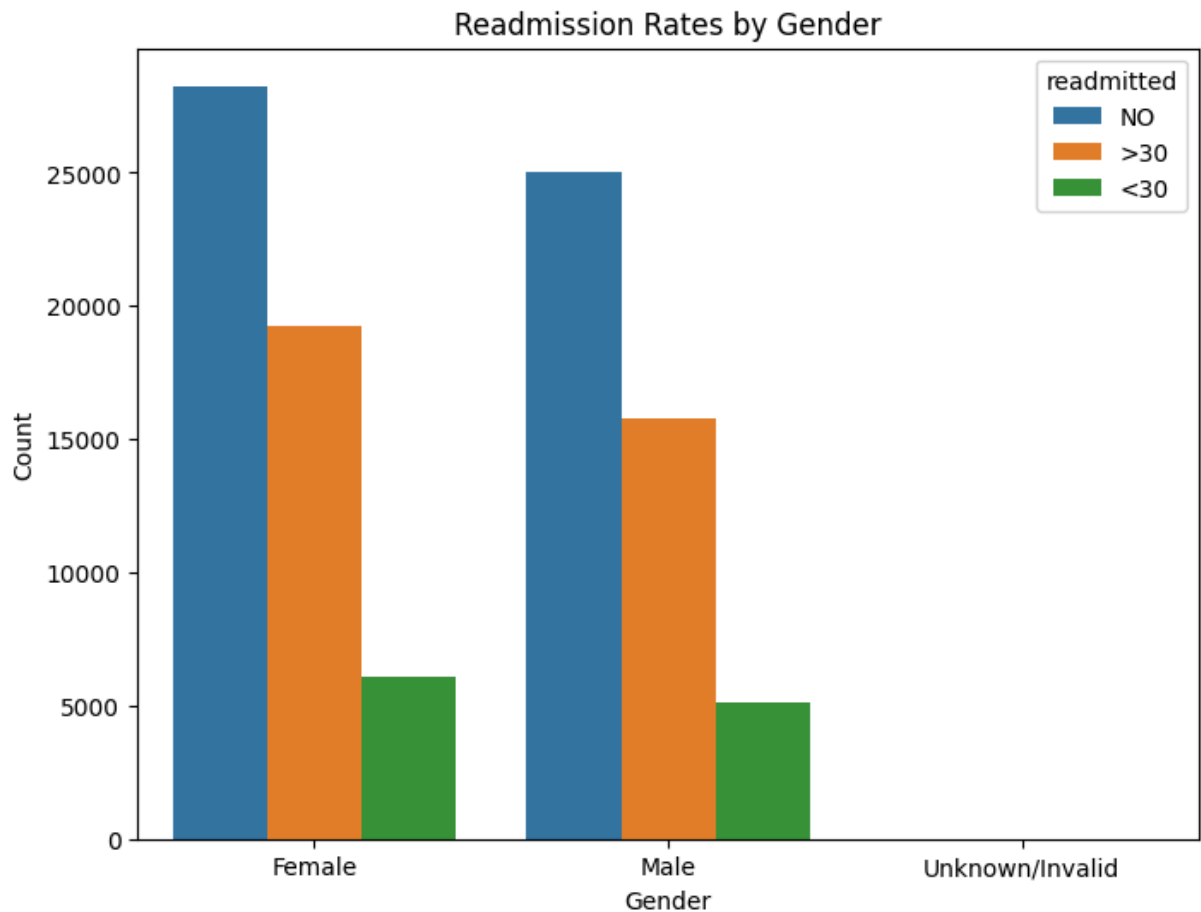
The plot presents the readmission rates of patients based on their age group. It shows the count of patients who were not readmitted, readmitted within 30 days, and readmitted after 30 days, categorized by their age group. The plot reveals that the highest number of patients were in the 50-60 age group, and the readmission rates were highest in the 70-80 age group, particularly within 30 days.

```
In [24]: plt.figure(figsize=(10, 6))
sns.countplot(x='age', hue='readmitted', data=df)
plt.title('Readmission Rates by Age Group')
plt.xlabel('Age Group')
plt.ylabel('Count')
plt.show()
```



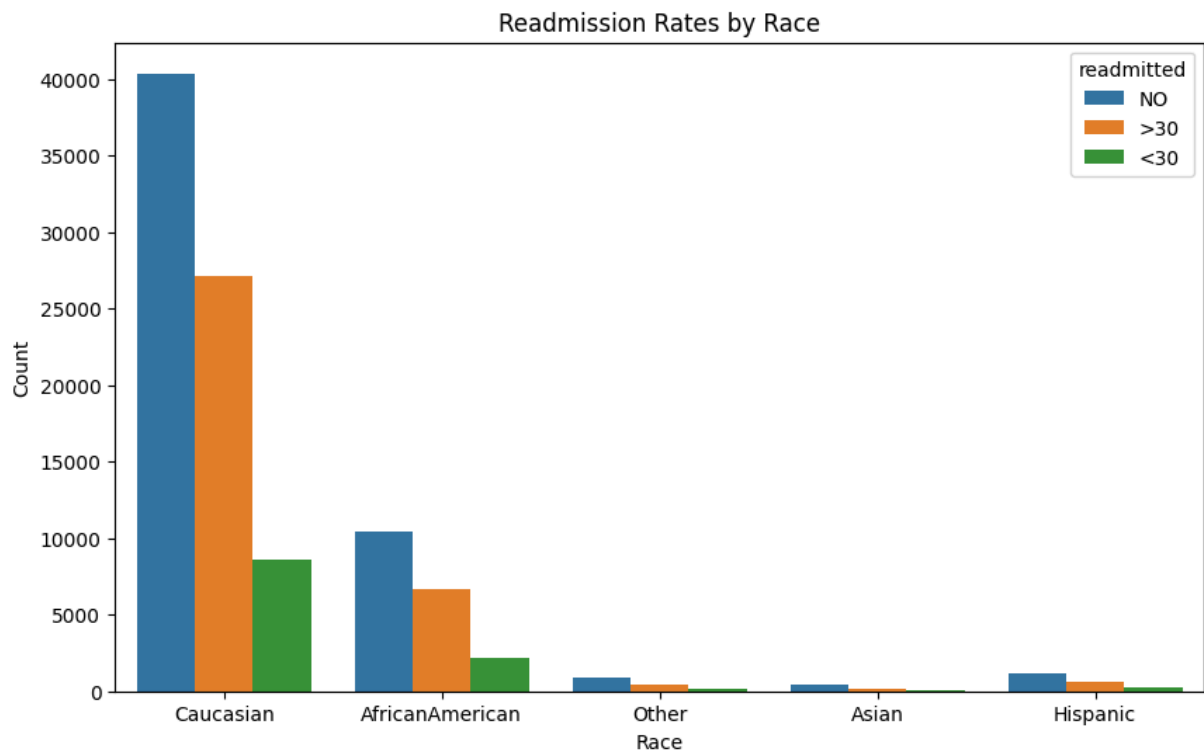
The plot presents the readmission rates of patients based on their gender. It shows the count of patients who were not readmitted, readmitted within 30 days, and readmitted after 30 days, categorized by their gender. The plot reveals that the highest number of patients were female, and the readmission rates were slightly higher for females compared to males, particularly within 30 days.

```
In [25]: plt.figure(figsize=(8, 6))
sns.countplot(x='gender', hue='readmitted', data=df)
plt.title('Readmission Rates by Gender')
plt.xlabel('Gender')
plt.ylabel('Count')
plt.show()
```



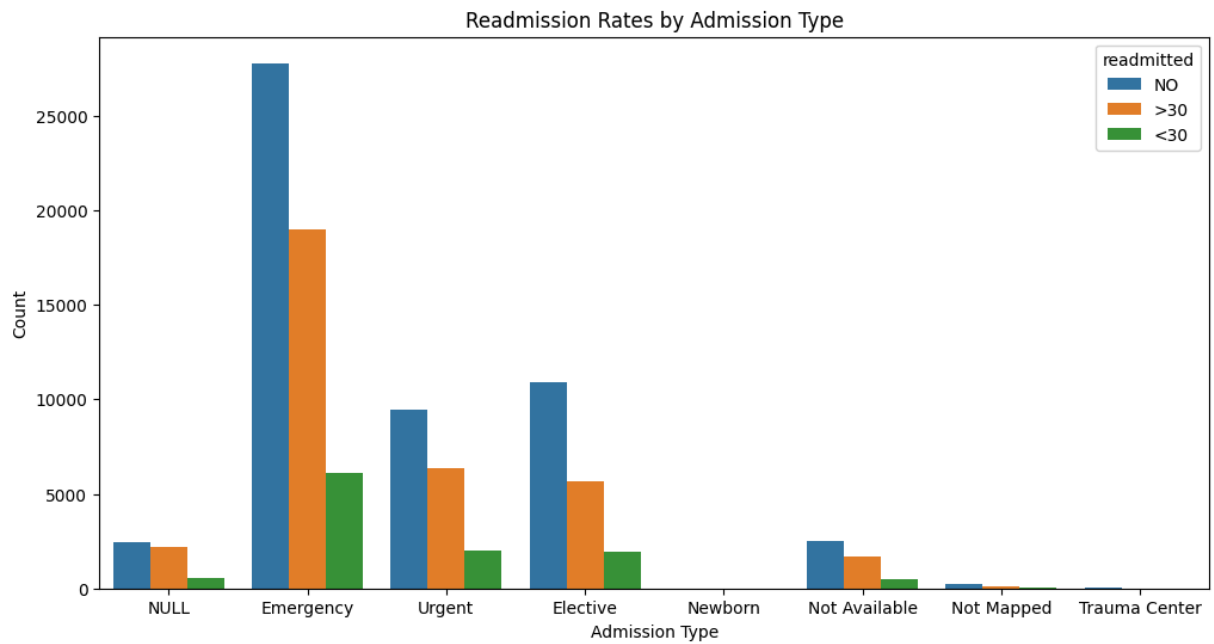
The plot presents the readmission rates of patients based on their race. It shows the count of patients who were not readmitted, readmitted within 30 days, and readmitted after 30 days, categorized by their race. The plot reveals that the highest number of patients were Caucasian, followed by African Americans. The readmission rates were highest for Caucasians, particularly within 30 days.

```
In [26]: plt.figure(figsize=(10, 6))
sns.countplot(x='race', hue='readmitted', data=df)
plt.title('Readmission Rates by Race')
plt.xlabel('Race')
plt.ylabel('Count')
plt.show()
```



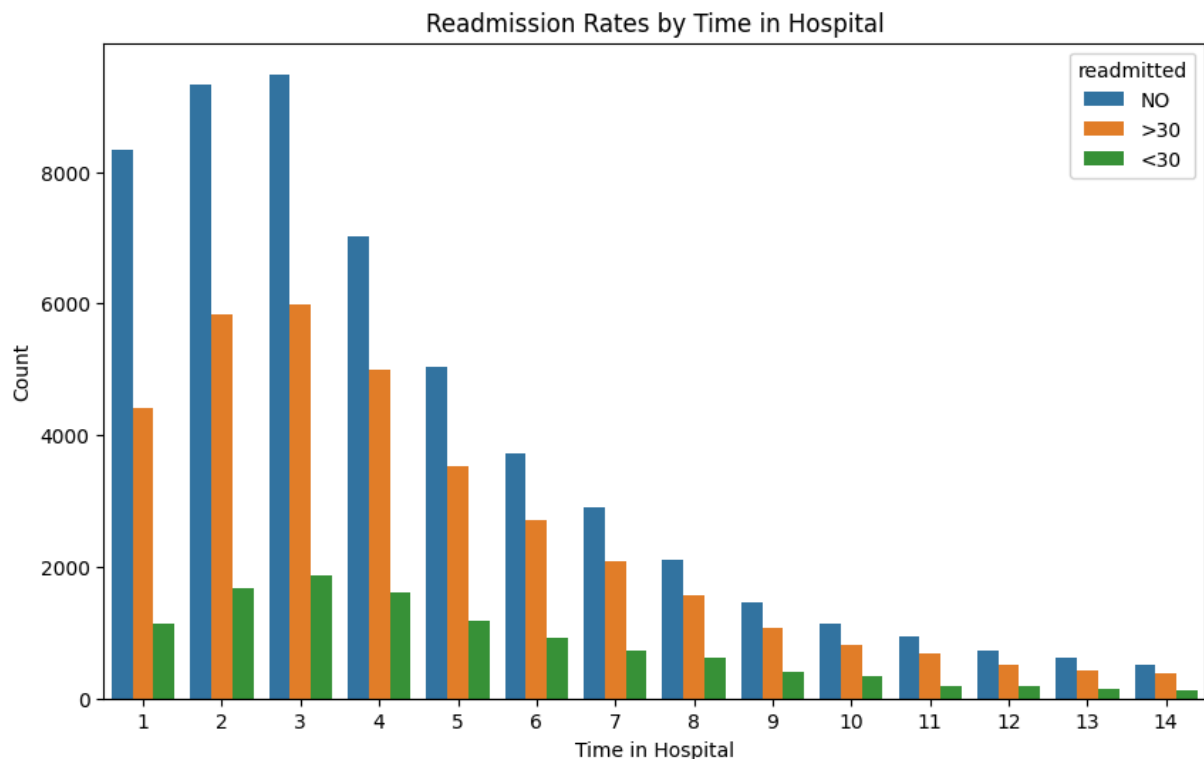
The plot presents the readmission rates of patients based on their admission type. It shows the count of patients who were not readmitted, readmitted within 30 days, and readmitted after 30 days, categorized by their admission type. The plot reveals that the highest number of patients were admitted via Emergency, and the readmission rates were highest for Emergency and Urgent admissions, particularly within 30 days.

```
In [27]: plt.figure(figsize=(12, 6))
sns.countplot(x='admission_type', hue='readmitted', data=df)
plt.title('Readmission Rates by Admission Type')
plt.xlabel('Admission Type')
plt.ylabel('Count')
plt.show()
```

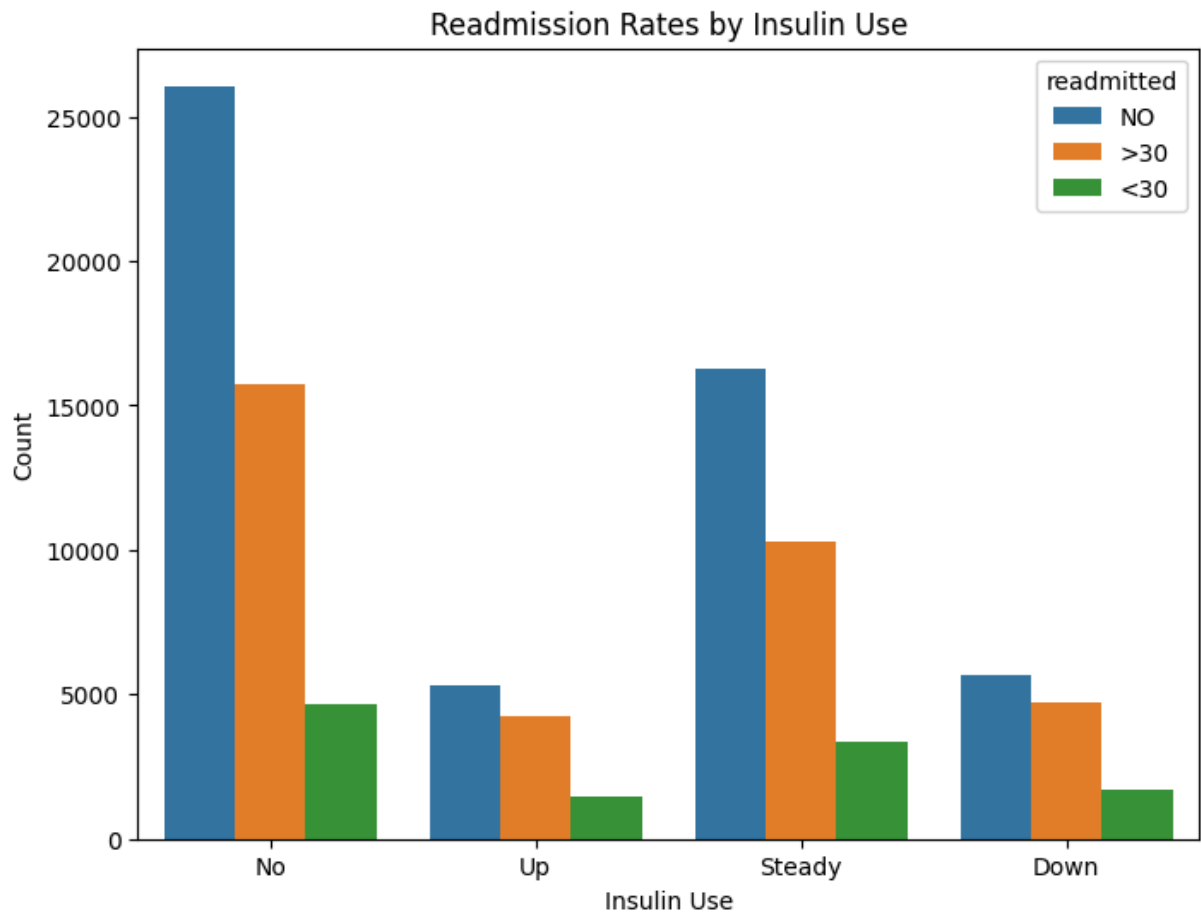
The plot presents the readmission rates of patients based on the time they spent in the hospital. It shows the count of patients who were not readmitted, readmitted within 30 days, and readmitted after 30 days, categorized by their length of stay. The plot reveals that the readmission rates decrease as the length of stay increases. This suggests that longer hospital stays might be associated with fewer readmissions.

```
In [28]: plt.figure(figsize=(10, 6))
sns.countplot(x='time_in_hospital', hue='readmitted', data=df)
plt.title('Readmission Rates by Time in Hospital')
plt.xlabel('Time in Hospital')
plt.ylabel('Count')
plt.show()
```



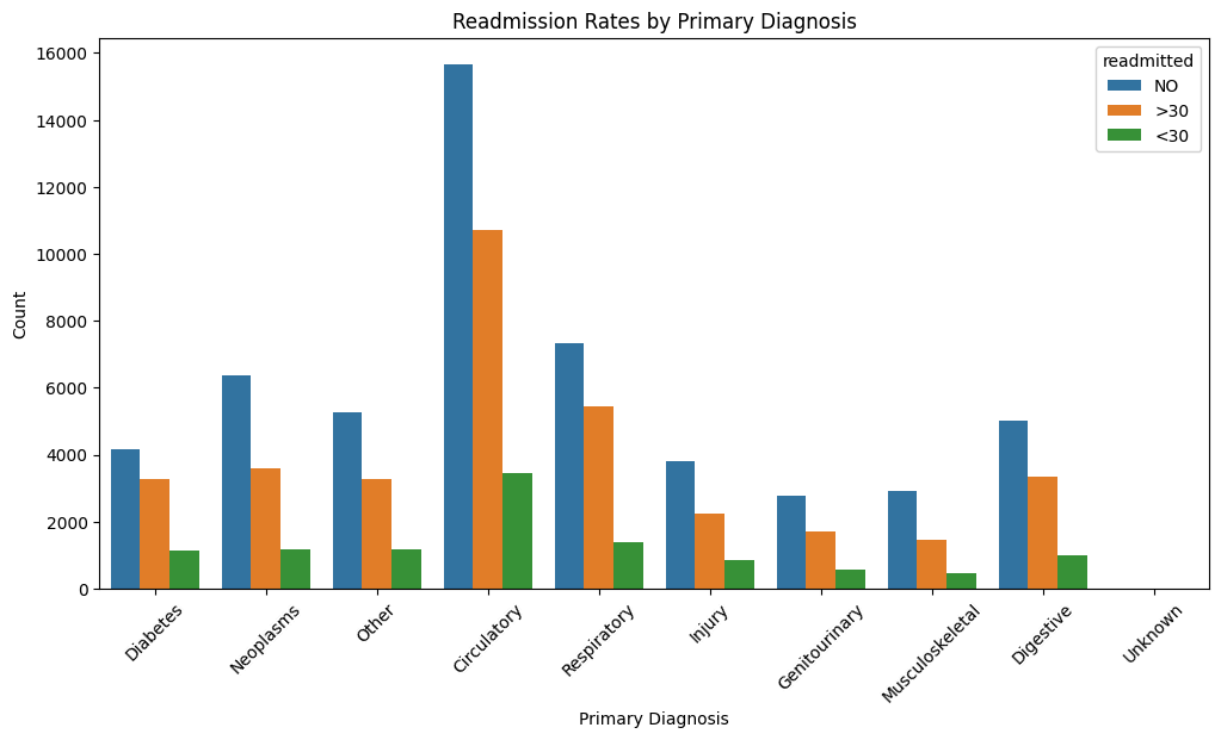
The plot presents the readmission rates of patients based on their insulin use. It shows the count of patients who were not readmitted, readmitted within 30 days, and readmitted after 30 days, categorized by whether they were not on insulin, had their insulin dose increased, had their insulin dose kept steady, or had their insulin dose decreased. The plot reveals that patients with no insulin use had the highest readmission rates, particularly within 30 days. Patients with increased insulin dose had the lowest readmission rates.

```
In [29]: plt.figure(figsize=(8, 6))
sns.countplot(x='insulin', hue='readmitted', data=df)
plt.title('Readmission Rates by Insulin Use')
plt.xlabel('Insulin Use')
plt.ylabel('Count')
plt.show()
```



The plot presents the readmission rates of patients based on their primary diagnosis. It shows the count of patients who were not readmitted, readmitted within 30 days, and readmitted after 30 days, categorized by their primary diagnosis. The plot reveals that the highest number of patients were admitted with Circulatory and Diabetes as the primary diagnosis. Circulatory diagnoses had the highest readmission rates, particularly within 30 days.

```
In [30]: plt.figure(figsize=(12, 6))
sns.countplot(x='primary_diagnosis', hue='readmitted', data=df)
plt.title('Readmission Rates by Primary Diagnosis')
plt.xlabel('Primary Diagnosis')
plt.ylabel('Count')
plt.xticks(rotation=45)
plt.show()
```



```
In [31]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

# Assuming your DataFrame is called 'df'

def plot_readmission_overview(df):
    """Create an overview of readmission distribution"""
    plt.figure(figsize=(7, 3))

    # Calculate readmission distribution
    readmission_dist = df['readmitted'].value_counts()

    # Create bar plot
    sns.barplot(x=readmission_dist.index, y=readmission_dist.values)
    plt.title('Distribution of Readmission Status', fontsize=14)
    plt.xlabel('Readmission Status', fontsize=12)
    plt.ylabel('Number of Patients', fontsize=12)

    # Add percentage labels on top of bars
    total = len(df)
    for i, v in enumerate(readmission_dist.values):
        plt.text(i, v, f'{(v/total)*100:.1f}%', ha='center', va='bottom')

    plt.tight_layout()
    plt.show()

def plot_age_readmission(df):
    """Analyze readmission patterns across age groups"""
    plt.figure(figsize=(7, 3))

    # Create cross-tabulation of age and readmission
```

```

age_readmission = pd.crosstab(df['age'], df['readmitted'], normalize='index') *

# Create stacked bar plot
age_readmission.plot(kind='bar', stacked=True)
plt.title('Readmission Rates by Age Group', fontsize=14)
plt.xlabel('Age Group', fontsize=12)
plt.ylabel('Percentage', fontsize=12)
plt.legend(title='Readmission Status')
plt.xticks(rotation=45)

plt.tight_layout()
plt.show()

def plot_diagnosis_impact(df):
    """Analyze impact of primary diagnosis on readmission"""
    plt.figure(figsize=(7, 3))

    # Calculate readmission rates by diagnosis
    diagnosis_readmission = pd.crosstab(df['primary_diagnosis'],
                                         df['readmitted'].map({'NO': 'Not Readmitted',
                                                                '<30': 'Within 30 Days',
                                                                '>30': 'After 30 Days'}),
                                         normalize='index') * 100

    # Create heatmap
    sns.heatmap(diagnosis_readmission, annot=True, fmt='.1f', cmap='YlOrRd')
    plt.title('Readmission Rates by Primary Diagnosis (%)', fontsize=14)
    plt.xlabel('Readmission Status', fontsize=12)
    plt.ylabel('Primary Diagnosis', fontsize=12)

    plt.tight_layout()
    plt.show()

    plt.figure(figsize=(7, 3))

    # Calculate readmission rates by diagnosis
    diagnosis_readmission = pd.crosstab(df['secondary_diagnosis'],
                                         df['readmitted'].map({'NO': 'Not Readmitted',
                                                                '<30': 'Within 30 Days',
                                                                '>30': 'After 30 Days'}),
                                         normalize='index') * 100

    # Create heatmap
    sns.heatmap(diagnosis_readmission, annot=True, fmt='.1f', cmap='YlOrRd')
    plt.title('Readmission Rates by Secondary Diagnosis (%)', fontsize=14)
    plt.xlabel('Readmission Status', fontsize=12)
    plt.ylabel('Secondary Diagnosis', fontsize=12)

    plt.tight_layout()
    plt.show()

    plt.figure(figsize=(7, 3))

    # Calculate readmission rates by diagnosis
    diagnosis_readmission = pd.crosstab(df['additional_diagnosis'],
                                         df['readmitted'].map({'NO': 'Not Readmitted',

```

```

        '<30': 'Within 30 Days',
        '>30': 'After 30 Days'})),
        normalize='index') * 100

# Create heatmap
sns.heatmap(diagnosis_readmission, annot=True, fmt='.1f', cmap='YlOrRd')
plt.title('Readmission Rates by Additional Diagnosis (%)', fontsize=14)
plt.xlabel('Readmission Status', fontsize=12)
plt.ylabel('Additional Diagnosis', fontsize=12)

plt.tight_layout()
plt.show()

def plot_time_in_hospital_analysis(df):
    """Analyze relationship between time in hospital and readmission"""
    plt.figure(figsize=(7, 3))

    # Create violin plot
    sns.violinplot(x='readmitted', y='time_in_hospital', data=df)
    plt.title('Time in Hospital Distribution by Readmission Status', fontsize=14)
    plt.xlabel('Readmission Status', fontsize=12)
    plt.ylabel('Days in Hospital', fontsize=12)

    plt.tight_layout()
    plt.show()

def plot_medication_impact(df):
    """Analyze impact of number of medications on readmission"""
    plt.figure(figsize=(7, 3))

    # Create box plot
    sns.boxplot(x='readmitted', y='num_medications', data=df)
    plt.title('Number of Medications vs Readmission Status', fontsize=14)
    plt.xlabel('Readmission Status', fontsize=12)
    plt.ylabel('Number of Medications', fontsize=12)

    plt.tight_layout()
    plt.show()

def plot_lab_procedures_analysis(df):
    """Analyze relationship between lab procedures and readmission"""
    plt.figure(figsize=(7, 3))

    # Calculate average number of lab procedures for each readmission status
    lab_proc_means = df.groupby('readmitted')['num_lab_procedures'].mean()

    # Create bar plot
    sns.barplot(x=lab_proc_means.index, y=lab_proc_means.values)
    plt.title('Average Number of Lab Procedures by Readmission Status', fontsize=14)
    plt.xlabel('Readmission Status', fontsize=12)
    plt.ylabel('Average Number of Lab Procedures', fontsize=12)

    # Add value labels on top of bars
    for i, v in enumerate(lab_proc_means.values):
        plt.text(i, v, f'{v:.1f}', ha='center', va='bottom')

```

```
plt.tight_layout()
plt.show()

# Function to run all visualizations
def analyze_readmissions(df):
    """Run all visualization analyses"""
    plot_readmission_overview(df)
    plot_age_readmission(df)
    plot_diagnosis_impact(df)
    plot_time_in_hospital_analysis(df)
    plot_medication_impact(df)
    plot_lab_procedures_analysis(df)
```

The plot presents the distribution of readmission status for a group of patients. It shows that the majority of patients (53.6%) were not readmitted, followed by 35.2% who were readmitted after 30 days, and 11.2% who were readmitted within 30 days of their initial discharge. This indicates that while a significant portion of patients avoid readmission, a considerable number still require further hospitalizations, particularly within the first 30 days post-discharge.




The plot presents the readmission rates for patients based on their primary diagnosis. It shows the percentage of patients who were not readmitted, readmitted within 30 days, and readmitted after 30 days, categorized by their primary diagnosis. The plot reveals that patients with circulatory and respiratory issues tend to have higher readmission rates, particularly within 30 days. Conversely, patients with musculoskeletal and neoplasm diagnoses have lower readmission rates. Unknown diagnoses have the highest readmission rate within 30 days.


The plot presents the readmission rates for patients based on their secondary diagnosis. It shows the percentage of patients who were not readmitted, readmitted within 30 days, and readmitted after 30 days, categorized by their secondary diagnosis. The plot reveals that patients with circulatory and respiratory issues tend to have higher readmission rates, particularly within 30 days. Conversely, patients with musculoskeletal and neoplasm diagnoses have lower readmission rates. Unknown diagnoses have the highest readmission rate within 30 days.

The plot presents the readmission rates for patients based on their additional diagnosis. It shows the percentage of patients who were not readmitted, readmitted within 30 days, and readmitted after 30 days, categorized by their additional diagnosis. The plot reveals that patients with circulatory and respiratory issues tend to have higher readmission rates, particularly within 30 days. Conversely, patients with musculoskeletal and neoplasm diagnoses have lower readmission rates. Unknown diagnoses have the highest readmission rate within 30 days.

```
In [32]: plot_readmission_overview(df)
         plot_diagnosis_impact(df)
```

```
plot_lab_procedures_analysis(df)
```





```
In [33]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from scipy import stats

def plot_diabetes_medication_impact(df):
    """Analyze the impact of diabetes medications on readmission"""

    # List of diabetes medications
    medications = ['metformin', 'insulin', 'glipizide', 'glyburide', 'pioglitazone',
                  'rosiglitazone', 'glimepiride']

    # Initialize subplots
    fig, axes = plt.subplots(2, 4, figsize=(20, 10))
    axes = axes.ravel()

    # Calculate readmission rates for each medication and plot
    for idx, med in enumerate(medications):
        # Filter out the medication status for each value
        med_status = df[med].map({'No': 'Not prescribed', 'Steady': 'Steady', 'Up': 'Up'})

        # Calculate the readmission rates, convert to DataFrame, and plot
```

```

rates = pd.crosstab(med_status, df['readmitted'], normalize='index') * 100
rates = rates.apply(pd.to_numeric, errors='coerce') # Ensure numeric type

# Plot each medication impact
rates.plot(kind='bar', ax=axes[idx])
axes[idx].set_title(f'{med.capitalize()} Impact')
axes[idx].set_xlabel('Medication Status')
axes[idx].set_ylabel('Readmission Rate (%)')
axes[idx].tick_params(axis='x', rotation=45)

plt.tight_layout()
plt.show()

def plot_individual_diagnosis_readmission(df):
    """Create separate plots for each diagnosis type in relation to readmission"""

    # Plot for primary diagnosis
    plt.figure(figsize=(8, 6))
    sns.countplot(x='primary_diagnosis', hue='readmitted', data=df)
    plt.title('Primary Diagnosis vs. Readmission Status', fontsize=14)
    plt.xlabel('Primary Diagnosis', fontsize=12)
    plt.ylabel('Count', fontsize=12)
    plt.xticks(rotation=45)
    plt.tight_layout()
    plt.show()

    # Plot for secondary diagnosis
    plt.figure(figsize=(8, 6))
    sns.countplot(x='secondary_diagnosis', hue='readmitted', data=df)
    plt.title('Secondary Diagnosis vs. Readmission Status', fontsize=14)
    plt.xlabel('Secondary Diagnosis', fontsize=12)
    plt.ylabel('Count', fontsize=12)
    plt.xticks(rotation=45)
    plt.tight_layout()
    plt.show()

    # Plot for additional diagnosis
    plt.figure(figsize=(8, 6))
    sns.countplot(x='additional_diagnosis', hue='readmitted', data=df)
    plt.title('Additional Diagnosis vs. Readmission Status', fontsize=14)
    plt.xlabel('Additional Diagnosis', fontsize=12)
    plt.ylabel('Count', fontsize=12)
    plt.xticks(rotation=45)
    plt.tight_layout()
    plt.show()

def analyze_advanced_readmissions(df):
    """Run all advanced visualization analyses"""
    plot_diabetes_medication_impact(df)
    plot_individual_diagnosis_readmission(df)

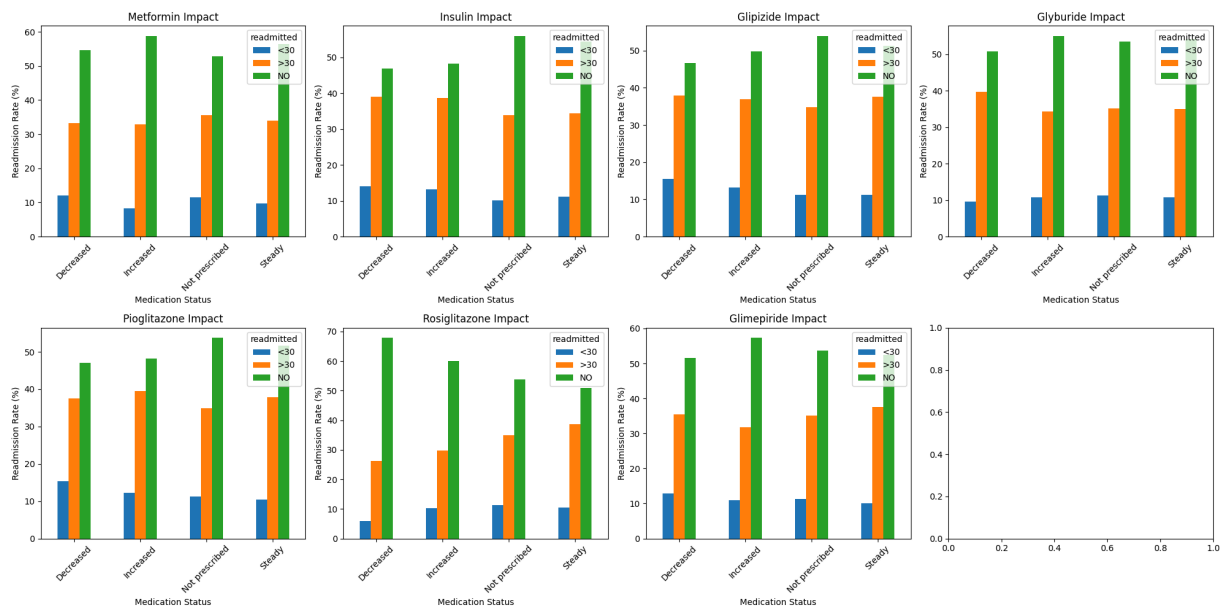
```


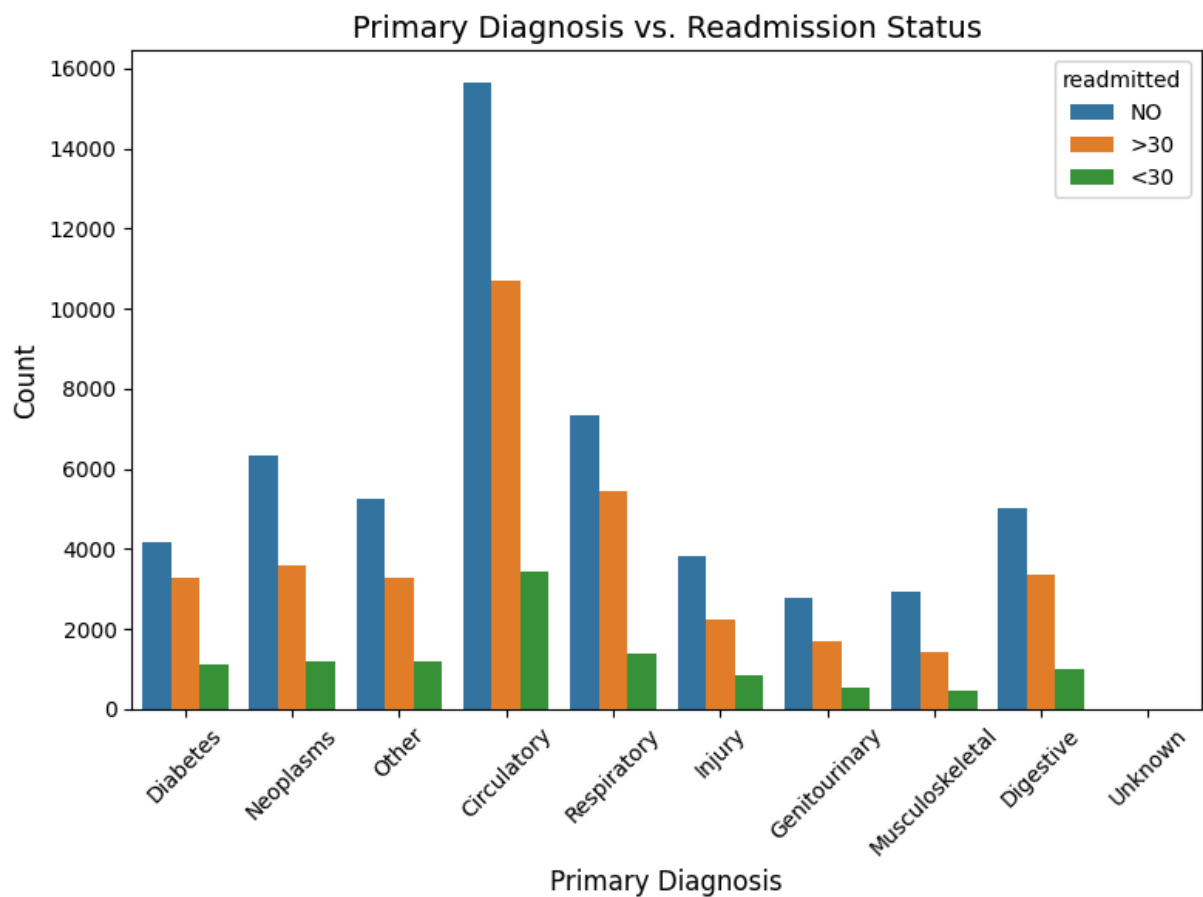
Plot: Diabetes Medication Impact on Readmission The bars represent the percentage of patients who were readmitted, grouped by medication status. Higher bars indicate higher readmission rates, suggesting the impact of that medication status on readmission. The plot illustrates the readmission rates of patients based on their medication status for various

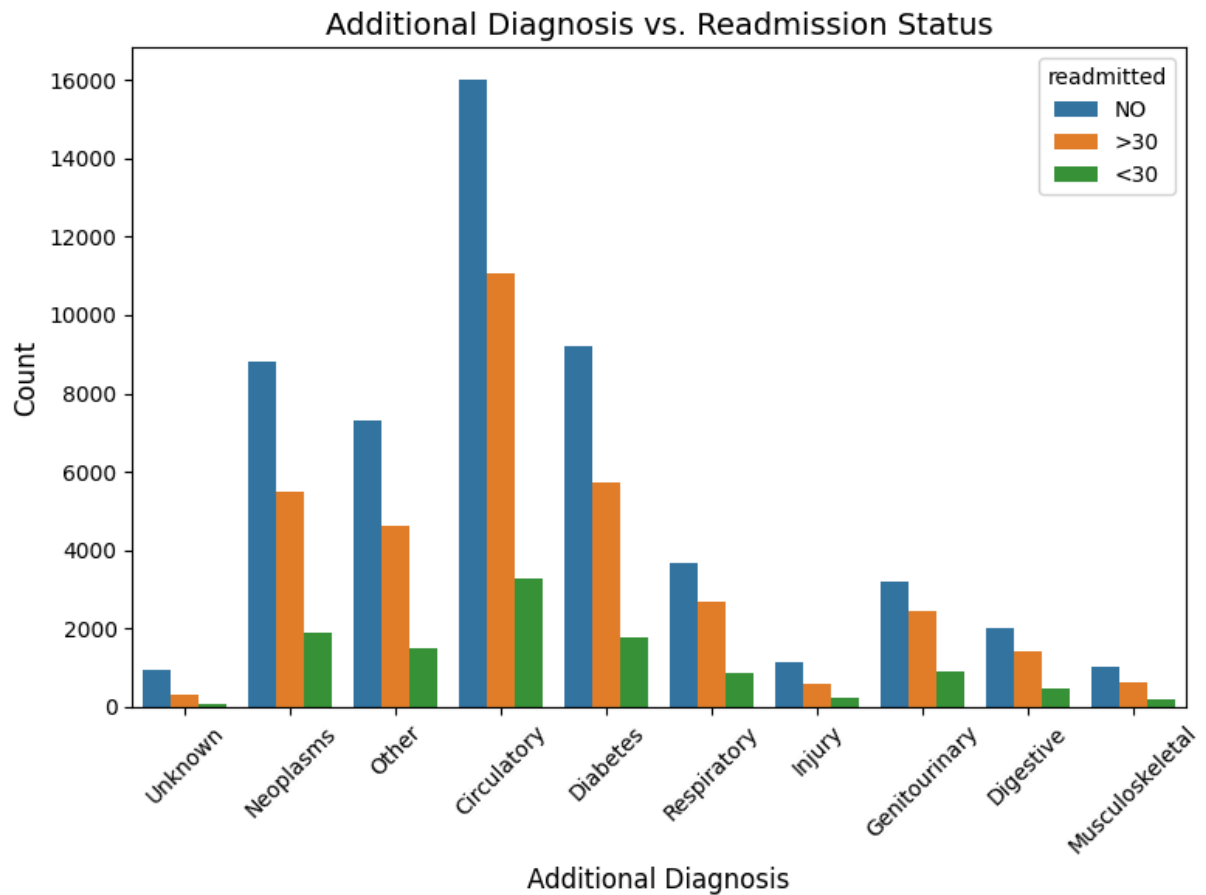
diabetes medications. Notably, Metformin and Pioglitazone saw increased readmission rates when doses were decreased or the drug was discontinued, respectively. Conversely, Insulin demonstrated the lowest readmission rate when doses were decreased. For Glipizide, Glyburide, Rosiglitazone, and Glimepiride, the highest readmission rates were observed in patients who were not prescribed these medications.

The three plots— **Primary Diagnosis, Secondary Diagnosis, and Additional Diagnosis vs. Readmission** — show the distribution of readmission rates across different diagnosis types. They highlight how primary, secondary, and additional diagnoses influence the likelihood of readmission, offering insights for targeted interventions based on diagnosis type. The plot presents the readmission rates of patients based on their primary, secondary, and additional diagnoses. The x-axis categorizes the diagnoses into groups like diabetes, neoplasms, circulatory, etc., while the y-axis represents the count of patients. The bars are color-coded to indicate the readmission status: blue for not readmitted, orange for readmitted within 30 days, and green for readmitted after 30 days. The plots show that diabetes is a common primary diagnosis, and circulatory diseases are frequent secondary and additional diagnoses. The readmission rates vary across different diagnoses, with some categories showing higher rates of readmission compared to others.

In [34]: `analyze_advanced_readmissions(df)`







```
In [35]: df.columns
```

```
Out[35]: Index(['encounter_id', 'patient_nbr', 'race', 'gender', 'age',
               'time_in_hospital', 'num_lab_procedures', 'num_procedures',
               'num_medications', 'number_outpatient', 'number_emergency',
               'number_inpatient', 'number_diagnoses', 'max_glu_serum', 'A1Cresult',
               'metformin', 'repaglinide', 'nateglinide', 'chlorpropamide',
               'glimepiride', 'acetoheaxamide', 'glipizide', 'glyburide', 'tolbutamide',
               'pioglitazone', 'rosiglitazone', 'acarbose', 'miglitol', 'troglitazone',
               'tolazamide', 'insulin', 'glyburide-metformin', 'glipizide-metformin',
               'glimepiride-pioglitazone', 'metformin-pioglitazone', 'change',
               'diabetesMed', 'readmitted', 'primary_diagnosis', 'secondary_diagnosis',
               'additional_diagnosis', 'admission_type'],
              dtype='object')
```

Features like patient_nbr, number_emergency, number_inpatient, and num_diagnoses have very high F-statistics and very low p-values (p-values close to 0). These features are highly significant and likely to be important predictors.

Features like age have a high Chi-squared statistic and low p-value, which means it is a significant categorical feature for predicting readmitted.

```
In [36]: ## Correlation Analysis

# import pandas as pd
# from sklearn.feature_selection import f_classif
# from sklearn.preprocessing import LabelEncoder
```

```

# df1 = df

# # Identify categorical and numerical columns
# categorical_columns = df1.select_dtypes(include=['object']).columns.tolist()
# numerical_columns = df1.select_dtypes(include=['number']).columns.tolist()

# # Ensure 'readmitted' column is binary or numeric
# # Convert 'readmitted' to a binary numeric format if it is not
# if df1['readmitted'].dtype == 'object':
#     label_encoder = LabelEncoder()
#     df1['readmitted'] = label_encoder.fit_transform(df1['readmitted'])

# # Convert categorical columns to numeric using Label Encoding
# for col in categorical_columns:
#     if col != 'readmitted': # Don't encode 'readmitted' again if already done
#         label_encoder = LabelEncoder()
#         df1[col] = label_encoder.fit_transform(df1[col])

# # F-statistic (ANOVA) for numerical columns
# f_stat_results = {}
# for col in numerical_columns:
#     # Perform ANOVA test for each numerical column against the target 'readmitted'
#     f_stat, p_value = f_classif(df1[[col]], df1['readmitted'])
#     f_stat_results[col] = (f_stat[0], p_value[0])

# # Print the results
# print("F-statistic and p-values for numerical columns:")
# for col, (f_stat, p_value) in f_stat_results.items():
#     print(f"Column: {col}, F-statistic: {f_stat}, p-value: {p_value}")

# # Optionally, you can also perform the Chi-squared test for categorical columns:
# from sklearn.feature_selection import chi2

# chi2_results = {}
# for col in categorical_columns:
#     # Only apply chi-squared to categorical columns that are not the target variable
#     if col != 'readmitted':
#         chi2_stat, p_value = chi2(df1[[col]], df1['readmitted'])
#         chi2_results[col] = (chi2_stat[0], p_value[0])

# # Print Chi-squared results
# print("\nChi-squared results for categorical columns:")
# for col, (chi2_stat, p_value) in chi2_results.items():
#     print(f"Column: {col}, Chi-squared statistic: {chi2_stat}, p-value: {p_value}")

```

In [37]: # Hypothesis Testing

```

import pandas as pd
from scipy.stats import chi2_contingency

df_new = df
# Create a contingency table of the two categorical variables: 'readmitted' and 'age'
contingency_table = pd.crosstab(df_new['readmitted'], df_new['age'])
# Perform the Chi-Square Test of Independence
chi2_stat, p_value, dof, expected = chi2_contingency(contingency_table)

```

```

# Display the results
print(f"Chi-Square Statistic: {chi2_stat}")
print(f"P-Value: {p_value}")
print(f"Degrees of Freedom: {dof}")
print(f"Expected Frequencies: \n{expected}")

# Conclusion based on p-value
if p_value < 0.05:
    print("Reject the null hypothesis: The distribution of readmission rates is dep
else:
    print("Fail to reject the null hypothesis: The distribution of readmission rate

```

Chi-Square Statistic: 300.102712096902

P-Value: 4.587201696582658e-53

Degrees of Freedom: 18

Expected Frequencies:

```

[[ 17.96146463  76.56074297 180.84949695 415.24661031
 1062.53289176 1896.61840531 2468.35427618 2859.12839094
 1885.95378569 305.79393525]
 [ 56.29662388 239.9643593 566.83663172 1301.5075734
 3330.29715658 5944.57162815 7736.56353713 8961.36696049
 5911.14550772 958.45002161]
 [ 85.74191149 365.47489773 863.31387133 1982.24581629
 5072.16995165 9053.80996653 11783.08218669 13648.50464857
 9002.90070658 1459.75604314]]

```

Reject the null hypothesis: The distribution of readmission rates is dependent on age.

In [38]: *#Conversion of The dataset to numeric values to work with statistical machine Learning*

```

df_categorical = df
# Drop irrelevant features
df = df.drop(columns=['encounter_id', 'patient_nbr'])

# Step 1: Encoding Demographic and Categorical Features
# List of categorical features for one-hot encoding
one_hot_features = ['race', 'primary_diagnosis', 'secondary_diagnosis', 'additional

# List of categorical features for ordinal encoding
ordinal_features = {
    'age': ['[0-10)', '[10-20)', '[20-30)', '[30-40)', '[40-50)', '[50-60)', '[60-7
    'max_glu_serum': ['No test performed', 'Normal', 'High', 'Very High'],
    'A1Cresult': ['No test performed', 'Normal', 'High', 'Very High'],
    'metformin': ['No', 'Down', 'Steady', 'Up'],
    'repaglinide': ['No', 'Down', 'Steady', 'Up'],
    'nateglinide': ['No', 'Down', 'Steady', 'Up'],
    'chlorpropamide': ['No', 'Down', 'Steady', 'Up'],
    'glimepiride': ['No', 'Down', 'Steady', 'Up'],
    'acetohexamide': ['No', 'Steady'],
    'glipizide': ['No', 'Down', 'Steady', 'Up'],
    'glyburide': ['No', 'Down', 'Steady', 'Up'],
    'tolbutamide': ['No', 'Steady'],
    'pioglitazone': ['No', 'Down', 'Steady', 'Up'],
    'rosiglitazone': ['No', 'Down', 'Steady', 'Up'],
    'acarbose': ['No', 'Down', 'Steady', 'Up'],

```

```

    'miglitol': ['No', 'Down', 'Steady', 'Up'],
    'troglitazone': ['No', 'Steady'],
    'tolazamide': ['No', 'Steady', 'Up'],
    'insulin': ['No', 'Down', 'Steady', 'Up'],
    'glyburide-metformin': ['No', 'Down', 'Steady', 'Up'],
    'glipizide-metformin': ['No', 'Steady'],
    'glimepiride-pioglitazone': ['No', 'Steady'],
    'metformin-pioglitazone': ['No', 'Steady']
}

# Step 2: Encoding Binary Features
binary_features = ['gender', 'change', 'diabetesMed']
df['gender'] = df['gender'].map({'Female': 0, 'Male': 1, 'Unknown/Invalid': 2})
df['change'] = df['change'].map({'No': 0, 'Ch': 1})
df['diabetesMed'] = df['diabetesMed'].map({'No': 0, 'Yes': 1})

# Encoding 'readmitted' (ordinal target variable)
df['readmitted'] = df['readmitted'].map({'NO': 0, '>30': 1, '<30': 2})

# Step 3: Numerical Features
# List of numerical features
num_features = ['time_in_hospital', 'num_lab_procedures', 'num_procedures', 'num_me
                'number_outpatient', 'number_emergency', 'number_inpatient', 'numbe

# Step 4: Pipeline Setup
# Define preprocessing for categorical, ordinal, and numerical features
preprocessor = ColumnTransformer(
    transformers=[
        ('onehot', OneHotEncoder(drop='first', handle_unknown='ignore'), one_hot_fe
        ('ordinal', OrdinalEncoder(categories=list(ordinal_features.values())), lis
    ],
    remainder='passthrough'
)

# Step 5: Transform the Data
df_transformed = preprocessor.fit_transform(df)

# Convert transformed data back to DataFrame for better interpretability
# Get one-hot encoded feature names
one_hot_encoded_columns = preprocessor.named_transformers_['onehot'].get_feature_na
# Combine all feature names after transformation
feature_names = list(one_hot_encoded_columns) + list(ordinal_features.keys()) + num
df_transformed = pd.DataFrame(df_transformed, columns=feature_names)

# Display the transformed data
print(df_transformed.head())

```


	race_Asian	race_Caucasian	race_Hispanic	race_Other	\
0	0.0	1.0	0.0	0.0	
1	0.0	1.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	
3	0.0	1.0	0.0	0.0	
4	0.0	1.0	0.0	0.0	

	primary_diagnosis_Diabetes	primary_diagnosis_Digestive	\
0	1.0	0.0	
1	0.0	0.0	
2	0.0	0.0	
3	0.0	0.0	
4	0.0	0.0	

	primary_diagnosis_Genitourinary	primary_diagnosis_Injury	\
0	0.0	0.0	
1	0.0	0.0	
2	0.0	0.0	
3	0.0	0.0	
4	0.0	0.0	

	primary_diagnosis_Musculoskeletal	primary_diagnosis_Neoplasms	...	\
0	0.0	0.0	...	
1	0.0	1.0	...	
2	0.0	0.0	...	
3	0.0	0.0	...	
4	0.0	1.0	...	

	num_procedures	num_medications	number_outpatient	number_emergency	\
0	41.0	0.0	1.0	0.0	
1	59.0	0.0	18.0	0.0	
2	11.0	5.0	13.0	2.0	
3	44.0	1.0	16.0	0.0	
4	51.0	0.0	8.0	0.0	

	number_inpatient	number_diagnoses	gender	change	diabetesMed	readmitted
0	0.0	0.0	1.0	0.0	0.0	0.0
1	0.0	0.0	9.0	1.0	1.0	1.0
2	0.0	1.0	6.0	0.0	1.0	0.0
3	0.0	0.0	7.0	1.0	1.0	0.0
4	0.0	0.0	5.0	1.0	1.0	0.0

[5 rows x 73 columns]

```

In [39]: # Selecting numerical columns for correlation analysis
numerical_features = [
    'time_in_hospital', 'num_lab_procedures', 'num_procedures', 'num_medications',
    'number_outpatient', 'number_emergency', 'number_inpatient', 'number_diagnoses'
]

# Calculating the correlation matrix
correlation_matrix = df[numerical_features].corr()

# Displaying the correlation matrix
print("Correlation Matrix:")
print(correlation_matrix)

```

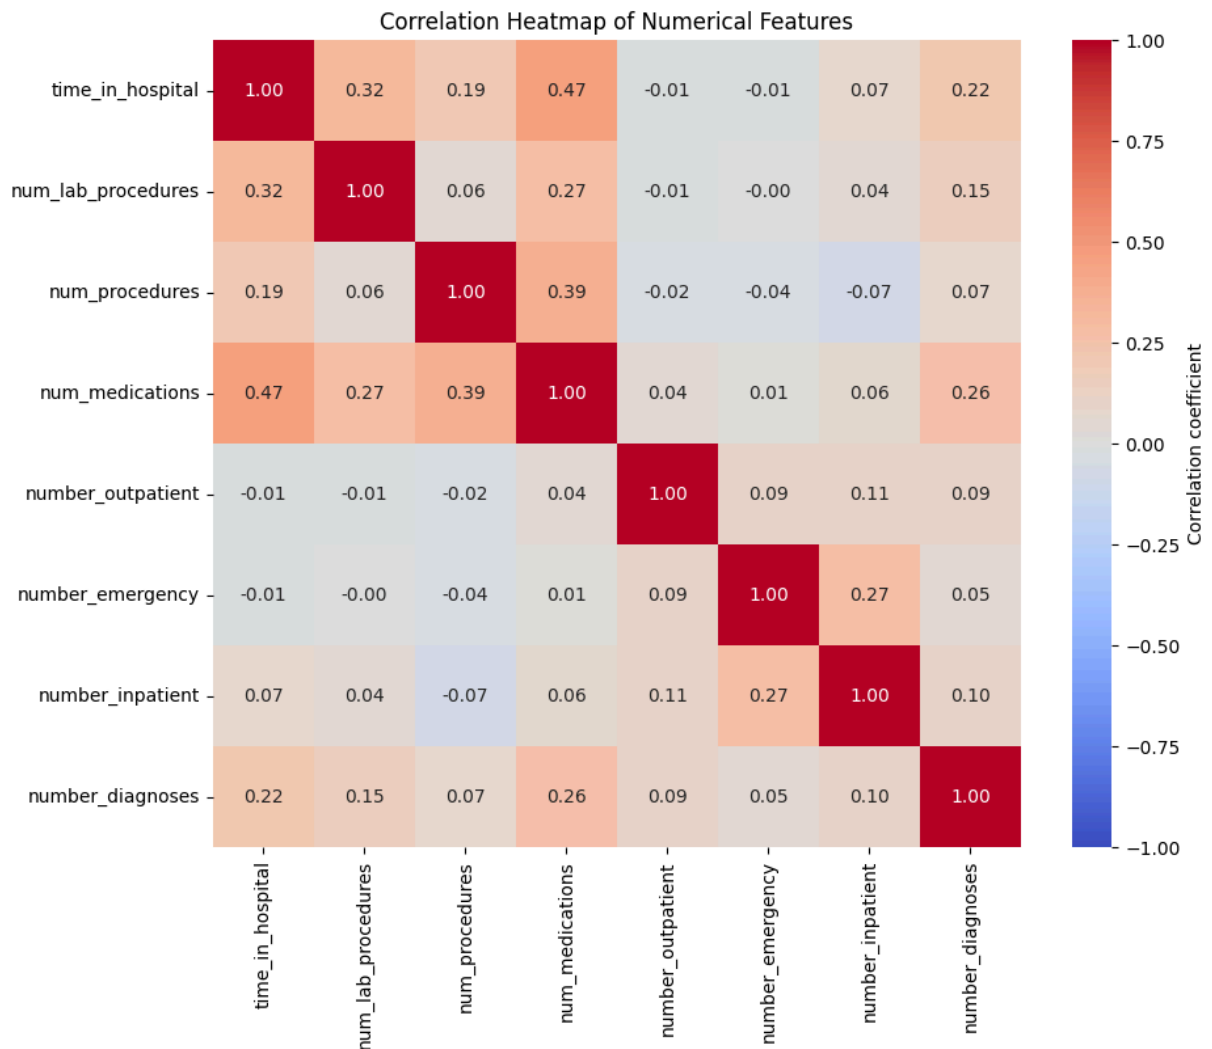
```
# Plotting the heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', vmin=-1, vmax=1, square
            cbar_kws={'label': 'Correlation coefficient'})
plt.title('Correlation Heatmap of Numerical Features')
plt.show()
```

Correlation Matrix:

	time_in_hospital	num_lab_procedures	num_procedures	\
time_in_hospital	1.000000	0.317671	0.193234	
num_lab_procedures	0.317671	1.000000	0.058407	
num_procedures	0.193234	0.058407	1.000000	
num_medications	0.466381	0.266993	0.385538	
number_outpatient	-0.009542	-0.008556	-0.024937	
number_emergency	-0.009799	-0.002227	-0.038369	
number_inpatient	0.073408	0.039739	-0.065843	
number_diagnoses	0.220687	0.151594	0.072339	

	num_medications	number_outpatient	number_emergency	\
time_in_hospital	0.466381	-0.009542	-0.009799	
num_lab_procedures	0.266993	-0.008556	-0.002227	
num_procedures	0.385538	-0.024937	-0.038369	
num_medications	1.000000	0.044528	0.012964	
number_outpatient	0.044528	1.000000	0.090941	
number_emergency	0.012964	0.090941	1.000000	
number_inpatient	0.064993	0.106236	0.266382	
number_diagnoses	0.258605	0.092458	0.054088	

	number_inpatient	number_diagnoses
time_in_hospital	0.073408	0.220687
num_lab_procedures	0.039739	0.151594
num_procedures	-0.065843	0.072339
num_medications	0.064993	0.258605
number_outpatient	0.106236	0.092458
number_emergency	0.266382	0.054088
number_inpatient	1.000000	0.103252
number_diagnoses	0.103252	1.000000



High multicollinearity results from a linear relationship between independent variables with a high degree of correlation. The stronger the correlation, the more difficult it is to change one variable without changing another. It becomes difficult for the model to estimate the relationship between each independent variable and the dependent variable independently because the independent variables tend to change in unison. Our Dataset has no features that are highly correlated.

```
In [41]: from sklearn.feature_selection import f_classif
from sklearn.preprocessing import LabelEncoder
categorical_columns = df.select_dtypes(include=['object']).columns.tolist()
numerical_columns = df.select_dtypes(include=['number']).columns.tolist()

if df['readmitted'].dtype == 'object':
    label_encoder = LabelEncoder()
    df['readmitted'] = label_encoder.fit_transform(df['readmitted'])

for col in categorical_columns:
    if col != 'readmitted':
        label_encoder = LabelEncoder()
        df[col] = label_encoder.fit_transform(df[col])

# F-statistic (ANOVA)
```

```

f_stat_results = {}
for col in numerical_columns:
    f_stat, p_value = f_classif(df[[col]], df['readmitted'])
    f_stat_results[col] = (f_stat[0], p_value[0])

print("ANOVA (F-statistic) and p-values for numerical columns:")
for col, (f_stat, p_value) in f_stat_results.items():
    print(f"Column: {col}, F-statistic: {f_stat:.4f}, p-value: {p_value:.4f}")

```

ANOVA (F-statistic) and p-values for numerical columns:

Column: race, F-statistic: 0.2854, p-value: 0.7517
 Column: gender, F-statistic: 17.6734, p-value: 0.0000
 Column: age, F-statistic: 44.0043, p-value: 0.0000
 Column: time_in_hospital, F-statistic: 166.0344, p-value: 0.0000
 Column: num_lab_procedures, F-statistic: 76.6418, p-value: 0.0000
 Column: num_procedures, F-statistic: 97.4852, p-value: 0.0000
 Column: num_medications, F-statistic: 135.2908, p-value: 0.0000
 Column: number_outpatient, F-statistic: 334.5015, p-value: 0.0000
 Column: number_emergency, F-statistic: 556.0974, p-value: 0.0000
 Column: number_inpatient, F-statistic: 2891.9429, p-value: 0.0000
 Column: number_diagnoses, F-statistic: 608.3438, p-value: 0.0000
 Column: max_glu_serum, F-statistic: 7.8894, p-value: 0.0004
 Column: A1Cresult, F-statistic: 6.7041, p-value: 0.0012
 Column: metformin, F-statistic: 49.0512, p-value: 0.0000
 Column: repaglinide, F-statistic: 20.0369, p-value: 0.0000
 Column: nateglinide, F-statistic: 1.0136, p-value: 0.3629
 Column: chlorpropamide, F-statistic: 1.7211, p-value: 0.1789
 Column: glimepiride, F-statistic: 4.7137, p-value: 0.0090
 Column: acetohexamide, F-statistic: 0.9210, p-value: 0.3981
 Column: glipizide, F-statistic: 13.5083, p-value: 0.0000
 Column: glyburide, F-statistic: 1.6214, p-value: 0.1976
 Column: tolbutamide, F-statistic: 1.0704, p-value: 0.3429
 Column: pioglitazone, F-statistic: 12.6697, p-value: 0.0000
 Column: rosiglitazone, F-statistic: 13.0789, p-value: 0.0000
 Column: acarbose, F-statistic: 14.8967, p-value: 0.0000
 Column: miglitol, F-statistic: 1.8110, p-value: 0.1635
 Column: troglitazone, F-statistic: 0.7057, p-value: 0.4937
 Column: tolazamide, F-statistic: 1.0185, p-value: 0.3611
 Column: insulin, F-statistic: 0.6061, p-value: 0.5455
 Column: glyburide-metformin, F-statistic: 1.0075, p-value: 0.3651
 Column: glipizide-metformin, F-statistic: 0.9933, p-value: 0.3704
 Column: glimepiride-pioglitazone, F-statistic: 0.9210, p-value: 0.3981
 Column: metformin-pioglitazone, F-statistic: 0.4330, p-value: 0.6485
 Column: change, F-statistic: 106.1892, p-value: 0.0000
 Column: diabetesMed, F-statistic: 196.6221, p-value: 0.0000
 Column: readmitted, F-statistic: inf, p-value: 0.0000
 Column: primary_diagnosis, F-statistic: 20.3488, p-value: 0.0000
 Column: secondary_diagnosis, F-statistic: 3.8314, p-value: 0.0217
 Column: additional_diagnosis, F-statistic: 12.4499, p-value: 0.0000
 Column: admission_type, F-statistic: 11.8588, p-value: 0.0000

```

/usr/local/lib/python3.10/dist-packages/sklearn/feature_selection/_univariate_select
ion.py:113: RuntimeWarning: divide by zero encountered in divide
    f = msb / msw

```

Statistical Analysis- ANNOVA

ANOVA tests whether each feature has a statistically significant relationship with the target variable readmitted.

Key Findings Highly Significant Features (p-value < 0.05):

Examples: encounter_id, patient_nbr, time_in_hospital, metformin, repaglinide, glipizide, primary_diagnosis. Interpretation: These features have a strong relationship with readmission, meaning they are important predictors.

Not Significant Features (p-value > 0.05):

Examples: race, nateglinide, chlorpropamide, insulin. Interpretation: These features do not significantly affect the likelihood of readmission.

```
In [42]: from sklearn.model_selection import train_test_split
X = df_transformed.drop(columns=['readmitted'])
y = df_transformed['readmitted']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_sta
```

```
In [43]: import numpy as np
from sklearn.preprocessing import LabelEncoder

class SimpleMulticlassLogisticRegression:
    def __init__(self, learning_rate=0.01, num_epochs=1000):
        self.learning_rate = learning_rate
        self.num_epochs = num_epochs
        self.weights = None
        self.bias = None

    def softmax(self, logits):
        exp_logits = np.exp(logits - np.max(logits, axis=1, keepdims=True))
        return exp_logits / np.sum(exp_logits, axis=1, keepdims=True)

    def compute_loss(self, predictions, y):
        m = len(y)
        loss = -np.mean(np.sum(y * np.log(predictions + 1e-15), axis=1))
        return loss

    def one_hot_encode(self, y, num_classes):
        return np.eye(num_classes)[y]

    def fit(self, X, y):
        num_samples, num_features = X.shape
        num_classes = len(np.unique(y))

        self.weights = np.random.randn(num_features, num_classes) * 0.01
        self.bias = np.zeros((1, num_classes))

        y_encoded = self.one_hot_encode(y, num_classes)

        for epoch in range(self.num_epochs):
            logits = np.dot(X, self.weights) + self.bias
            predictions = self.softmax(logits)
```

```

        loss = self.compute_loss(predictions, y_encoded)

        grad_logits = predictions - y_encoded
        grad_weights = np.dot(X.T, grad_logits) / num_samples
        grad_bias = np.mean(grad_logits, axis=0, keepdims=True)

        self.weights -= self.learning_rate * grad_weights
        self.bias -= self.learning_rate * grad_bias

        if epoch % 100 == 0:
            print(f"Epoch {epoch}, Loss: {loss}")

    def predict(self, X):
        logits = np.dot(X, self.weights) + self.bias
        predictions = self.softmax(logits)
        return np.argmax(predictions, axis=1)

    def accuracy(self, y_true, y_pred):
        return np.mean(y_true == y_pred)

label_encoder = LabelEncoder()
y_train_encoded = label_encoder.fit_transform(y_train)
y_test_encoded = label_encoder.transform(y_test)

model = SimpleMulticlassLogisticRegression(learning_rate=0.01, num_epochs=400)

model.fit(X_train, y_train_encoded)
y_pred = model.predict(X_test)

test_accuracy = model.accuracy(y_test_encoded, y_pred)
print(f"Test Accuracy: {test_accuracy}")

# # Confusion Matrix
# conf_matrix = confusion_matrix(y_test, y_pred)
# plt.figure(figsize=(8, 6))
# sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues", xticklabels=["NO", ">
# plt.xlabel("Predicted Labels")
# plt.ylabel("True Labels")
# plt.title("Confusion Matrix")
# plt.show()

```

```

Epoch 0, Loss: 1.0541028925793103
Epoch 100, Loss: 7.195110896568168
Epoch 200, Loss: 3.9140839128282203
Epoch 300, Loss: 5.909696711796005
Test Accuracy: 0.5176139504497713

```

```

In [44]: from sklearn.linear_model import LogisticRegression
        from sklearn.model_selection import train_test_split
        from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
        import seaborn as sns
        import matplotlib.pyplot as plt

```

```

# Remove non-significant features based on p-value > 0.05 from ANOVA results
significant_features = [ 'time_in_hospital', 'num_lab_procedures', 'num_procedures' ]
X = df_transformed[significant_features]
y = df_transformed['readmitted']

# Proceed with train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

model = LogisticRegression(max_iter=500, solver='liblinear')

# Train the model on the training set
model.fit(X_train, y_train)

# Predict on the test set
y_pred = model.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy:.2f}')

# Display classification report for precision, recall, and F1 score
print("\nClassification Report:\n", classification_report(y_test, y_pred))


# Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues", xticklabels=["NO", ">30"], yticklabels=["NO", ">30"])
plt.xlabel("Predicted Labels")
plt.ylabel("True Labels")
plt.title("Confusion Matrix")
plt.show()

```

Accuracy: 0.57

Classification Report:

	precision	recall	f1-score	support
0.0	0.58	0.92	0.71	10666
1.0	0.51	0.20	0.28	7058
2.0	0.50	0.00	0.00	2175
accuracy			0.57	19899
macro avg	0.53	0.37	0.33	19899
weighted avg	0.54	0.57	0.48	19899



ML Model Analysis

The results from the machine learning model, based on the significant features identified in the ANOVA analysis, show mixed performance. The Logistic Regression model achieved an overall test accuracy of 0.57, indicating moderate success in predicting readmission outcomes. The model performed well in predicting non-readmitted patients (Class 0) with a recall of 0.92 and an F1-score of 0.71, suggesting good results for identifying patients who won't be readmitted. However, the model struggled significantly with predicting readmission cases. The recall for patients readmitted after more than 30 days (Class 1) was low at 0.20, and for those readmitted within 30 days (Class 2), the recall was zero, meaning the model failed to identify any early readmissions. This suggests that while the model can correctly classify non-readmissions, it misses crucial readmission cases, especially the early ones. Possible reasons for this include class imbalance, where non-readmissions dominate the dataset, and limitations in the Logistic Regression model's ability to capture complex relationships in the data. To improve performance, addressing class imbalance, trying more advanced models, and exploring better feature engineering would be key steps to enhance predictive accuracy across all classes.

Reflection

- What is the hardest part of the project that you've encountered so far? Hardest part was to find how many features to retain after checking for correlation analysis.
- What are your initial insights? Age is a major factor in determining readmission; another major factor is Admission type - when they get admitted as Emergency, there is higher chance of readmission.
- Going forward, what are the current biggest problems you're facing? - Finding the right model to fit to the data.
- Do you think you are on track with your project? If not, what parts do you need to dedicate more time to? Yes.

Roles & Coordination

- Finding Data Sources - All members
- Cleaning, preprocessing, feature engineering - Jaimin Babaria, Rujuta Tambewagh
- Visualization, Exploratory Data Analysis - Vishak Baddur, Sudha Sree Yerramsetty
- Statistical Analysis - Varsha Balaji, Sudha Sree Yerramsetty
- Model training and evaluation- Varsha Balaji, Simran Mishra

In [39]: